

Chapter 2

Application Layer

Introduction: 1-2

2

Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



Application Layer: 2-3

3

Application layer: overview

Our goals:

- conceptual *and* implementation aspects of application-layer protocols
 - transport-layer service models
 - client-server paradigm
 - peer-to-peer paradigm
- learn about protocols by examining popular application-layer protocols and infrastructure
 - HTTP
 - SMTP, IMAP
 - DNS
 - video streaming systems, CDNs
- programming network applications
 - socket API

Application Layer: 2-4

4

Some network apps

- social networking
 - Web
 - text messaging
 - e-mail
 - multi-user network games
 - streaming stored video (YouTube, Hulu, Netflix)
 - P2P file sharing
 - voice over IP (e.g., Skype)
 - real-time video conferencing (e.g., Zoom)
 - Internet search
 - remote login
 - ...
- Q: your favorites?***

Application Layer: 2-5

5

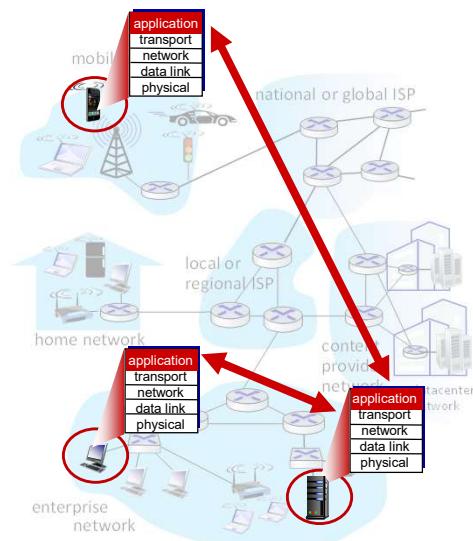
Creating a network app

write programs that:

- run on (different) end systems
- communicate over network
- e.g., web server software communicates with browser software

no need to write software for network-core devices

- network-core devices do not run user applications
- **applications on end systems** allows for **rapid app development, propagation**



Application Layer: 2-6

6

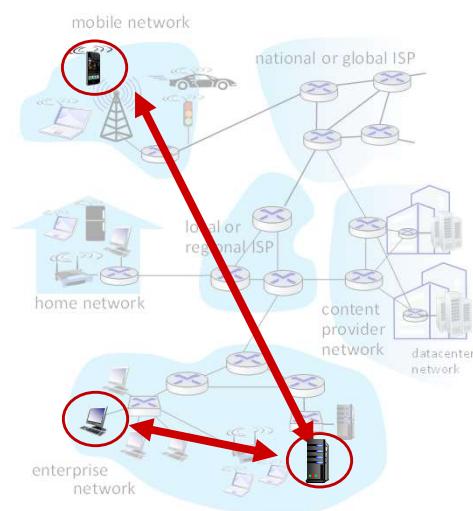
Client-server paradigm

server:

- always-on host
- permanent IP address
- often in data centers, for scaling

clients:

- contact, communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do *not* communicate directly with each other
- examples: HTTP, IMAP, FTP

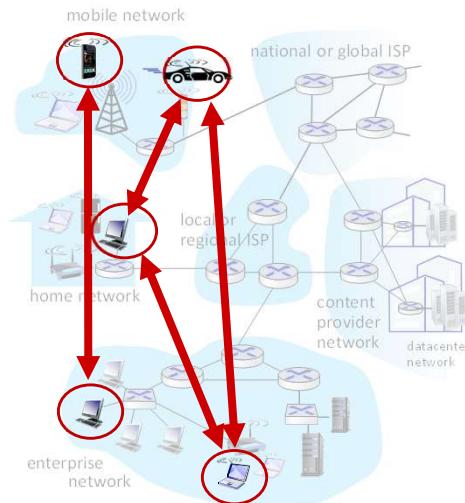


Application Layer: 2-7

7

Peer-peer architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
 - complex management
- example: P2P file sharing



Application Layer: 2-8

8

Processes communicating

- process:* program running within a host
- within same host, two processes communicate using **inter-process communication** (defined by OS)
 - processes in different hosts communicate by exchanging **messages**

clients, servers

client process: process that initiates communication

server process: process that waits to be contacted

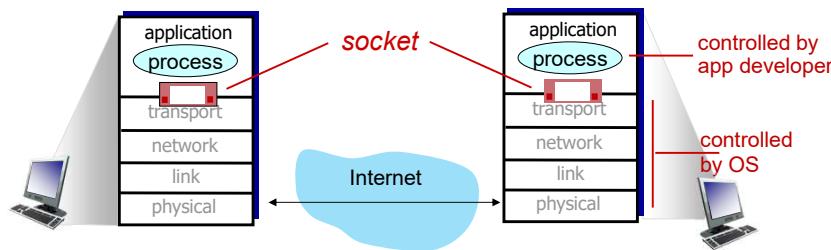
- note: applications with P2P architectures have client processes & server processes

Application Layer: 2-9

9

Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process
 - two sockets involved: one on each side



Application Layer: 2-10

10

Addressing processes

- to receive messages, process must have **identifier**
- host device has unique 32-bit IP address
- **Q:** does IP address of host on which process runs suffice for identifying the process?
 - **A:** no, *many* processes can be running on same host
- **identifier** includes both **IP address** and **port numbers** associated with process on host.
- example port numbers:
 - HTTP server: 80
 - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
 - **IP address:** 128.119.245.12
 - **port number:** 80
- more shortly...

Application Layer: 2-11

11

An application-layer protocol defines:

- types of messages exchanged,
 - e.g., request, response
- message syntax:
 - what fields in messages & how fields are delineated
- message semantics
 - meaning of information in fields
- rules for when and how processes send & respond to messages

open protocols:

- defined in RFCs, everyone has access to protocol definition
- allows for interoperability
- e.g., HTTP, SMTP

proprietary protocols:

- e.g., Skype, Zoom

Application Layer: 2-12

12

What transport service does an app need?

data integrity

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- other apps (“elastic apps”) make use of whatever throughput they get

security

- encryption, data integrity, ...

Application Layer: 2-13

13

Transport service requirements: common apps

application	data loss	throughput	time sensitive?
file transfer/download	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5Kbps-1Mbps video:10Kbps-5Mbps	yes, 10's msec
streaming audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	Kbps+	yes, 10's msec
text messaging	no loss	elastic	yes and no

Application Layer: 2-14

14

Internet transport protocols services

TCP service:

- **reliable transport** between sending and receiving process
- **flow control**: sender won't overwhelm receiver
- **congestion control**: throttle sender when network overloaded
- **connection-oriented**: setup required between client and server processes
- **does not provide**: timing, minimum throughput guarantee, security

UDP service:

- **unreliable data transfer** between sending and receiving process
- **does not provide**: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup.

Q: why bother? **Why** is there a UDP?

Application Layer: 2-15

15

Internet applications, and transport protocols

application	application layer protocol	transport protocol
file transfer/download	FTP [RFC 959]	TCP
e-mail	SMTP [RFC 5321]	TCP
Web documents	HTTP 1.1 [RFC 7320]	TCP
Internet telephony	SIP [RFC 3261], RTP [RFC 3550], or proprietary	TCP or UDP
streaming audio/video	HTTP [RFC 7320], DASH	TCP
interactive games	WOW, FPS (proprietary)	UDP or TCP

Application Layer: 2-16

16

Securing TCP

Vanilla TCP & UDP sockets:

- no encryption
- cleartext passwords sent into socket traverse Internet in cleartext (!)

Transport Layer Security (TLS)

- provides encrypted TCP connections
- data integrity
- end-point authentication

TSL implemented in application layer

- apps use TSL libraries, that use TCP in turn
- cleartext sent into “socket” traverse Internet *encrypted*
- more: Chapter 8

Application Layer: 2-17

17

Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



Application Layer: 2-18

18

Web and HTTP

First, a quick review...

- web page consists of **objects**, each of which can be stored on different Web servers
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of **base HTML-file** which includes **several referenced objects, each** addressable by a **URL**, e.g.,

`www.someschool.edu/someDept/pic.gif`

host name

path name

Application Layer: 2-19

19

Google Chrome, Internet Explorer, Opera, Firefox, Safari

HTTP overview

HTTP: hypertext transfer protocol

- Web's application-layer protocol
- client/server model:
 - *client*: browser that **requests**, **receives**, (using HTTP protocol) and **"displays"** Web objects
 - *server*: Web server sends (using HTTP protocol) objects in **response to requests**

Apache web server, Apache Tomcat, Microsoft's Internet Information Services (IIS) Windows Server, Nginx web server, lighttpd, Jigsaw, Klone,, Abyss web server, Oracle Web Tier, X5 (Xitami) web server, Zeus web server



Application Layer: 2-20

20

HTTP overview (continued)

HTTP uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between **browser** (HTTP client) and **Web server** (HTTP server)
- TCP connection closed

HTTP is “stateless”

- server maintains **no** information about past client requests

aside
protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client **crashes**, their views of “state” may be **inconsistent**, **must be reconciled**

Application Layer: 2-21

21

HTTP connections: two types

Non-persistent HTTP

1. TCP connection opened
 2. at most one object sent over TCP connection
 3. TCP connection closed
- downloading multiple objects required multiple connections

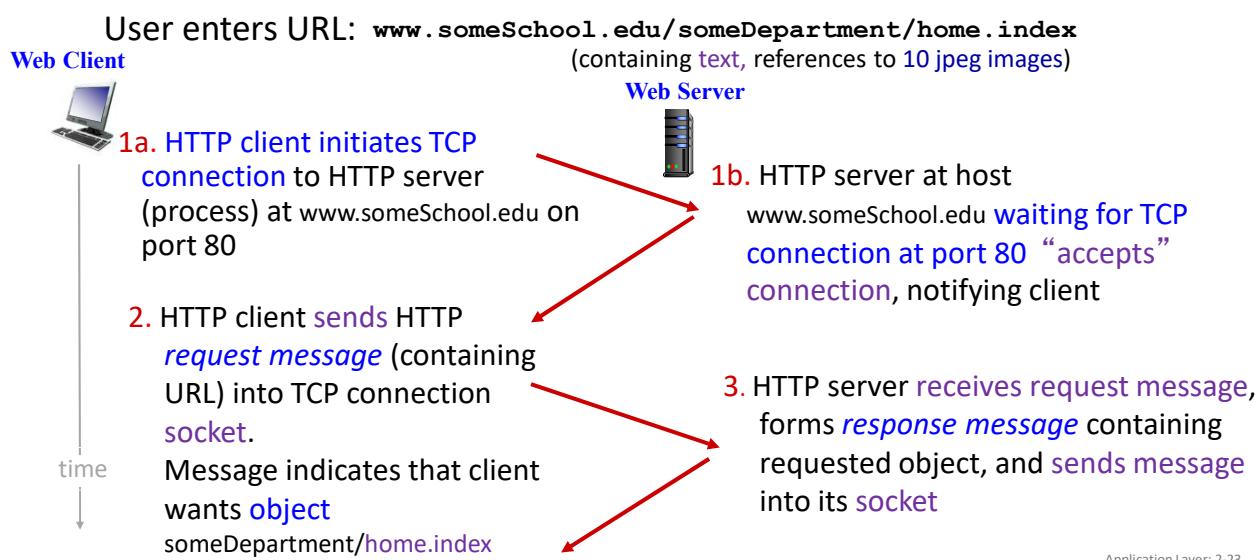
Persistent HTTP

- TCP connection opened to a server
- multiple objects can be sent over *single TCP connection* between client, and that server
- TCP connection closed

Application Layer: 2-22

22

Non-persistent HTTP: example

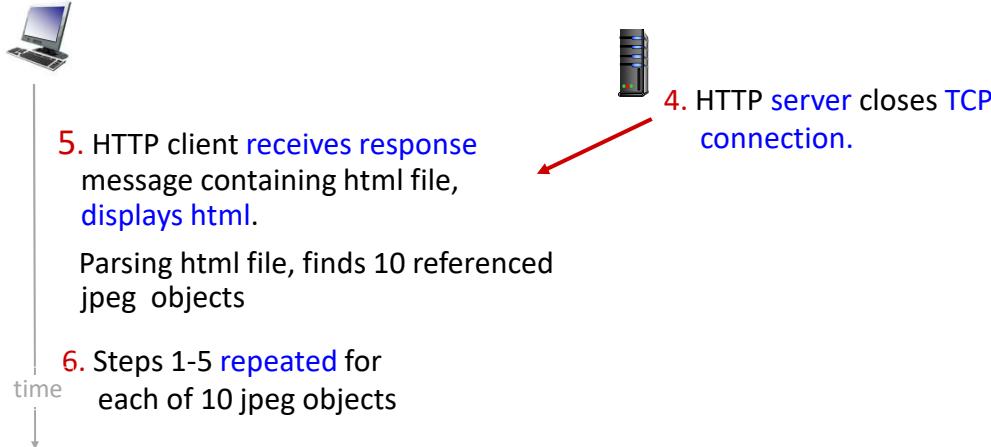


Application Layer: 2-23

23

Non-persistent HTTP: example (cont.)

User enters URL: `www.someSchool.edu/someDepartment/home.index`
 (containing text, references to 10 jpeg images)



Application Layer: 2-24

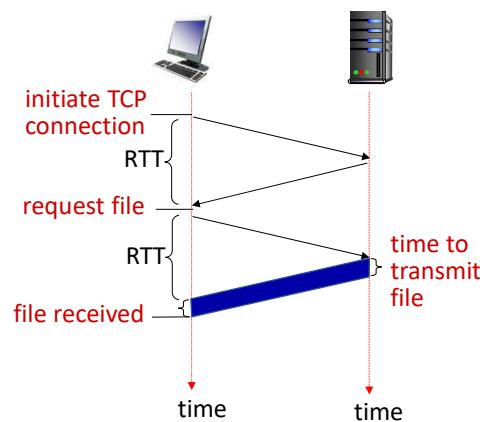
24

Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time (per object):

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- object/file transmission time



$$\text{Non-persistent HTTP response time} = 2\text{RTT} + \text{file transmission time}$$

Application Layer: 2-25

25

Persistent HTTP (HTTP 1.1)

Non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open multiple parallel TCP connections to fetch referenced objects in parallel

Persistent HTTP (HTTP1.1):

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects (cutting response time in half)

Application Layer: 2-26

26

Persistent HTTP: response time

Application Layer: 2-27

27

HTTP request message

- two types of HTTP messages: *request, response*
- **HTTP request message:**
 - ASCII (human-readable format)

request line (GET, POST,
HEAD commands) →

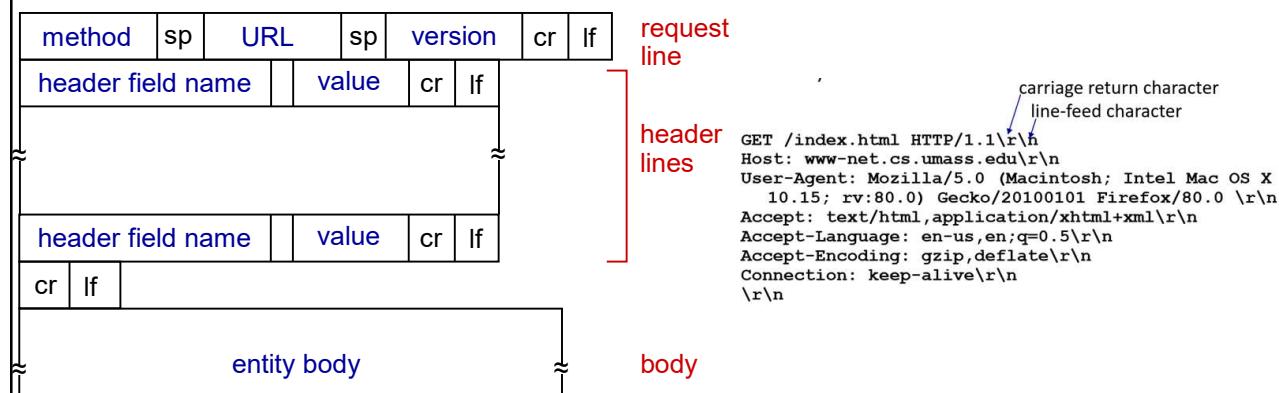
carriage return character
line-feed character

carriage return, line feed →
at start of line indicates
end of header lines

Application Layer: 2-28

28

HTTP request message: general format



Application Layer: 2-31

31

Other HTTP request messages

POST method:

- web page often includes **form input**
- **user input** sent from client to server in **entity body** of HTTP POST request message

GET method (for sending data to server):

- include user data in URL field of HTTP GET request message (following a '?'):

`www.somesite.com/animalsearch?monkeys&banana`

HEAD method:

- **requests headers** (only) that would be returned if specified URL were requested with an HTTP GET method.

PUT method:

- uploads new file (**object**) to server
- completely replaces file that exists at **specified URL** with content in **entity body** of POST HTTP request message

Application Layer: 2-32

32

HTTP response message

status line (protocol
status code status phrase)

→ HTTP/1.1 200 OK

header
lines

```
Date: Tue, 08 Sep 2020 00:53:20 GMT
Server: Apache/2.4.6 (CentOS)
        OpenSSL/1.0.2k-fips PHP/7.4.9
        mod_perl/2.0.11 Perl/v5.16.3
Last-Modified: Tue, 01 Mar 2016 18:57:50 GMT
ETag: "a5b-52d015789ee9e"
Accept-Ranges: bytes
Content-Length: 2651
Content-Type: text/html; charset=UTF-8
\r\n
```

data, e.g., requested
HTML file

→ data data data data data ...



Application Layer: 2-31

33

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (in Location: field)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

Application Layer: 2-36

36

Trying out HTTP (client side) for yourself

1. netcat to your favorite Web server:

```
% nc -c -v gaia.cs.umass.edu 80
```

- opens TCP connection to port 80 (default HTTP server port) at gaia.cs.umass.edu.
- anything typed in will be sent to port 80 at gaia.cs.umass.edu

2. type in a GET HTTP request:

```
GET /kurose_ross/interactive/index.php HTTP/1.1
Host: gaia.cs.umass.edu
```

- by typing this in (hit carriage return twice), you send this minimal (but complete) GET request to HTTP server

3. look at response message sent by HTTP server!

(or use Wireshark to look at captured HTTP request/response)

Application Layer: 2-37

37

Trying out HTTP (client side) for yourself

1. Telnet to your favorite Web server:

```
telnet cis.poly.edu 80
```

Opens TCP connection to port 80
(default HTTP server port) at cis.poly.edu.
Anything typed in sent
to port 80 at cis.poly.edu

2. Type in a GET HTTP request:

```
GET /~ross/ HTTP/1.1
Host: cis.poly.edu
```

By typing this in (hit carriage
return twice), you send
this minimal (but complete)
GET request to HTTP server

3. Look at response message sent by HTTP server!

2: Application Layer 38

38

The screenshot shows a Windows Telnet window titled 'Telnet unix.kmit.ac.th'. The user has typed 'telnet cis.poly.edu 80' and connected to the server. They then entered the HTTP request 'GET /~ross/ HTTP/1.1' followed by 'Host: cis.poly.edu'. The server responded with an HTTP/1.0 200 OK status and various headers. The response body contains HTML code for Keith Ross's homepage, including his title as Professor of Computer Science, his department, address, phone number, fax number, email, and bio.

```

$ telnet cis.poly.edu 80
Trying 128.238.32.126...
Connected to cis.poly.edu.
Escape character is '^J'.
GET /~ross/ HTTP/1.1
Host: cis.poly.edu

HTTP/1.0 200 OK
Date: Thu, 07 Jul 2005 17:50:13 GMT
Server: Apache/1.2.5
Last-Modified: Thu, 20 Jan 2005 19:19:17 GMT
ETag: "15834-2748-41f00435"
Content-Length: 10056
Accept-Ranges: bytes
Content-Type: text/html
X-Cache: MISS from proxy.net.kmit.ac.th
Connection: close

<html>
<head>
<title>Keith Ross's Homepage</title>
</head>
<body>

<br>&nbsp;
<h2>
Keith W. Ross</h2>

<h3> Leonard J. Shustek Professor of Computer Science</h3>
<p>
Department of Computer and Information Science
<br>Polytechnic University
<br>Six MetroTech Center
<br>Brooklyn, NY 11201
<br>Phone: 718-260-3859
<br>Fax: 718-260-3609
<p>e-mail:&nbsp; <a href="mailto:ross@poly.edu">ross@poly.edu</a>
<h3>
BioSketch</h3>
Professor Ross is the <a href="http://www.poly.edu/news/Keith_Ross_news.cfm">Chair Professor in Computer Science</a> at Polytechnic University. Before joining Polytechnic University, he was a professor in the Multimedia Communications Department at Eurecom Institute (1996-2003) and was a professor in

```

response message sent by HTTP server

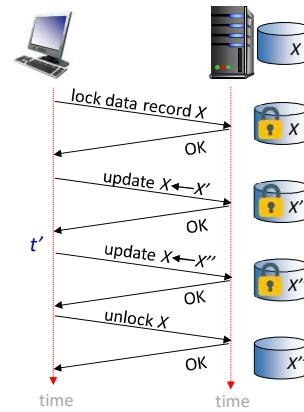
39

Maintaining user/server state: cookies

Recall: HTTP GET/response interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web “transaction”
 - no need for client/server to track “state” of multi-step exchange
 - all HTTP requests are independent of each other
 - no need for client/server to “recover” from a partially-completed-but-never-completely-completed transaction

a **stateful protocol**: client makes two changes to X, or none at all



Q: what happens if network connection or client crashes at t' ?

Application Layer: 2-40

40

Maintaining user/server state: cookies

Web sites and client browser use *cookies (RFC 6265)* to maintain some *state* between transactions

four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in *next* HTTP *request* message
- 3) cookie file kept on *user's host*, managed by user's browser
- 4) back-end database at Web site

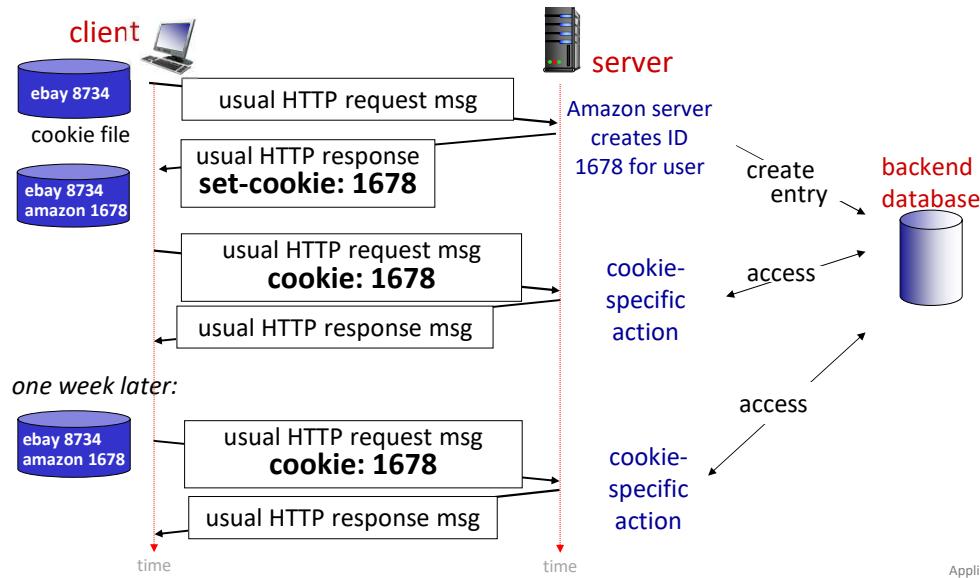
Example:

- Susan uses browser on laptop, visits specific e-commerce site for first time
- when initial HTTP request arrives at site, site creates:
 - unique ID (aka “cookie”)
 - entry in backend database for ID
- subsequent HTTP requests from Susan to this site will contain cookie ID value, allowing site to “identify” Susan

Application Layer: 2-41

41

Maintaining user/server state: cookies



42

HTTP cookies: comments

What cookies can be used for:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

Challenge: How to keep state?

- *at protocol endpoints:* maintain state at sender/receiver over multiple transactions
- *in messages:* cookies in HTTP messages carry state

aside
cookies and privacy:

- cookies permit sites to *learn* a lot about you on their site.
- third party persistent cookies (tracking cookies) allow common identity (cookie value) to be tracked across multiple web sites

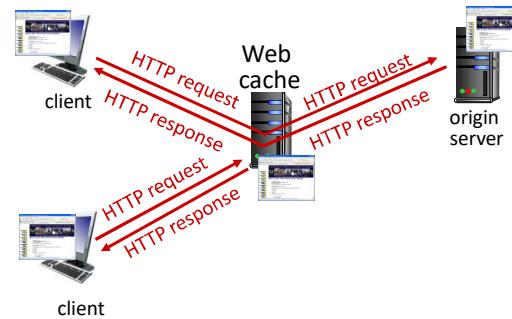
Application Layer: 2-43

43

Web caches

Goal: satisfy client requests **without involving origin server**

- user configures browser to point to a (local) **Web cache**
- browser sends all HTTP requests to cache
 - *if* object in cache: cache returns object to client
 - *else* cache requests object from origin server, caches received object, then returns object to client



Application Layer: 2-44

44

Web caches (aka proxy servers)

- Web cache acts as **both client and server**
 - server for original requesting client
 - client to origin server
- server tells cache about object's allowable caching in **response header**:

`Cache-Control: max-age=<seconds>`

`Cache-Control: no-cache`

Why Web caching?

- **reduce response time for client request**
 - cache is closer to client
- **reduce traffic** on an institution's access link
- **Internet is dense with caches**
 - enables “poor” content providers to more effectively deliver content

Application Layer: 2-45

45

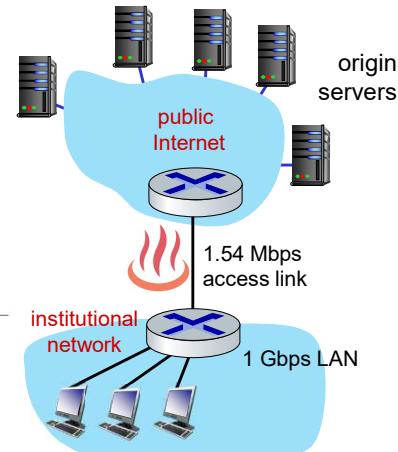
Caching example

Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps

Performance:

- access link utilization = **.97** *problem: large queueing delays at high utilization!*
- LAN utilization: .0015
- end-end delay = Internet delay +
access link delay + LAN delay
= 2 sec + **minutes** + usecs



Application Layer: 2-46

46

Option 1: buy a faster access link

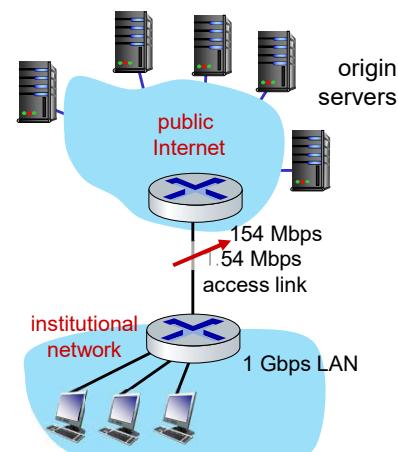
Scenario:

- access link rate: **154 Mbps**
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps

Performance:

- access link utilization = **.07** → .0097
- LAN utilization: .0015
- end-end delay = Internet delay +
access link delay + LAN delay
= 2 sec + **msecs** + usecs

Cost: faster access link (expensive!)



Application Layer: 2-47

47

Option 2: install a web cache

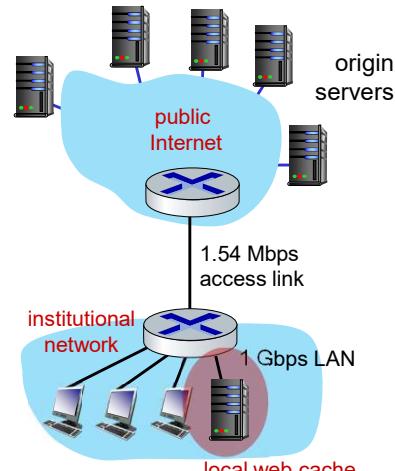
Scenario:

- access link rate: 1.54 Mbps
- RTT from institutional router to server: 2 sec
- web object size: 100K bits
- average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps

Cost: web cache (cheap!)

Performance:

- LAN utilization: .? *How to compute link utilization, delay?*
- access link utilization = ? *utilization, delay?*
- average end-end delay = ?



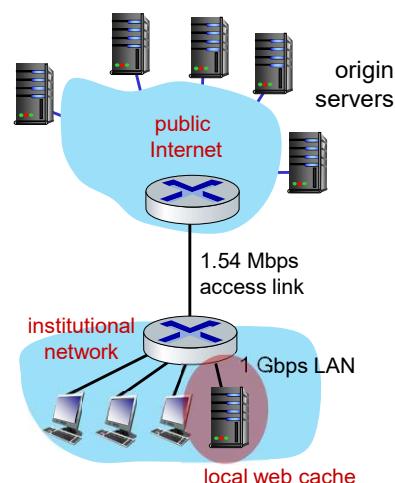
48

Calculating access link utilization, end-end delay with cache:

suppose cache **hit rate** is 0.4:

- 40% requests served by cache, with low (msec) delay
- 60% requests satisfied at origin
 - rate to browsers over **access link**
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
 - **access link utilization** = $0.9 / 1.54 = .58$ means low (msec) queueing delay at access link (10 msec.)
- average end-end delay:
 $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
 $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$

lower average end-end delay than with 154 Mbps link (and cheaper too!)

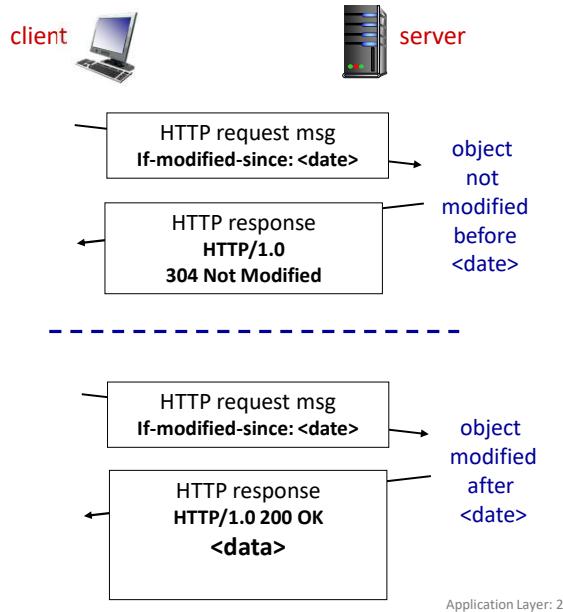


49

Conditional GET

Goal: don't send object if cache has up-to-date cached version

- no object transmission delay (or use of network resources)
- **client:** specify date of cached copy in HTTP request
If-modified-since: <date>
- **server:** response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified



Application Layer: 2-50

50

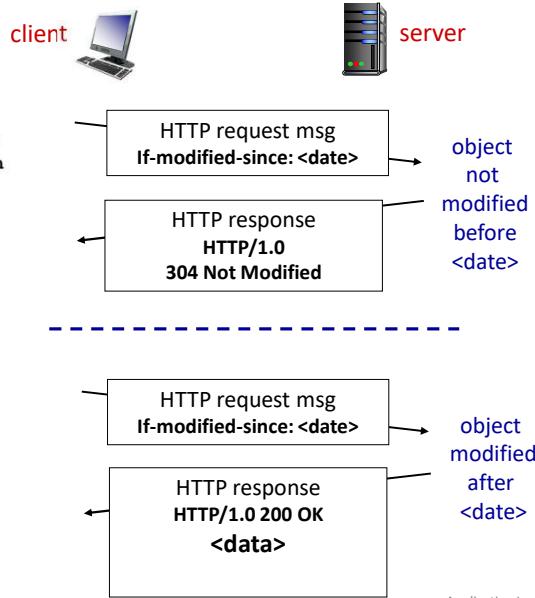
Conditional GET

```

GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X
  10.15; rv:80.0) Gecko/20100101 Firefox/80.0 \r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Connection: keep-alive\r\n
\r\n
HTTP/1.1 200 OK
Date: Tue, 08 Sep 2020 00:53:20 GMT
Server: Apache/2.4.6 (CentOS)
  OpenSSL/1.0.2k-fips PHP/7.4.9
  mod_perl/2.0.11 Perl/v5.16.3
Last-Modified: Tue, 01 Mar 2016 18:57:50 GMT
ETag: "a5b-52d015789ee9e"
Accept-Ranges: bytes
Content-Length: 2651
Content-Type: text/html; charset=UTF-8
\r\n
  data data data data ...

```

carriage return character
line-feed character



Application Layer: 2-51

51

HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP1.1: introduced multiple, pipelined GETs over single TCP connection

- server responds *in-order* (**FCFS**: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (**head-of-line (HOL) blocking**) behind large object(s)
- **loss recovery** (retransmitting lost TCP segments) stalls object transmission

Application Layer: 2-52

52

HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] increased flexibility at *server* in sending objects to client:

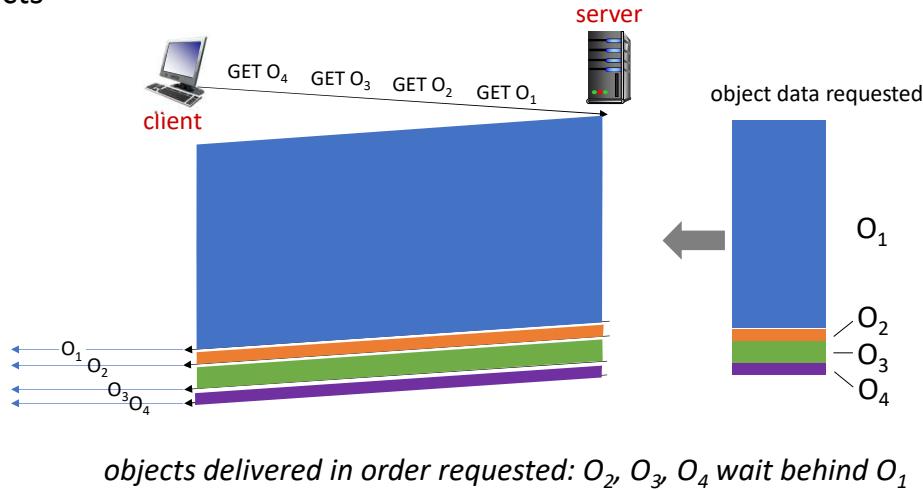
- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified **object priority** (not necessarily FCFS)
- **push** unrequested objects to client
- divide objects into **frames, schedule frames** to mitigate HOL blocking

Application Layer: 2-53

53

HTTP/2: mitigating HOL blocking

HTTP 1.1: client requests 1 large object (e.g., video file) and 3 smaller objects

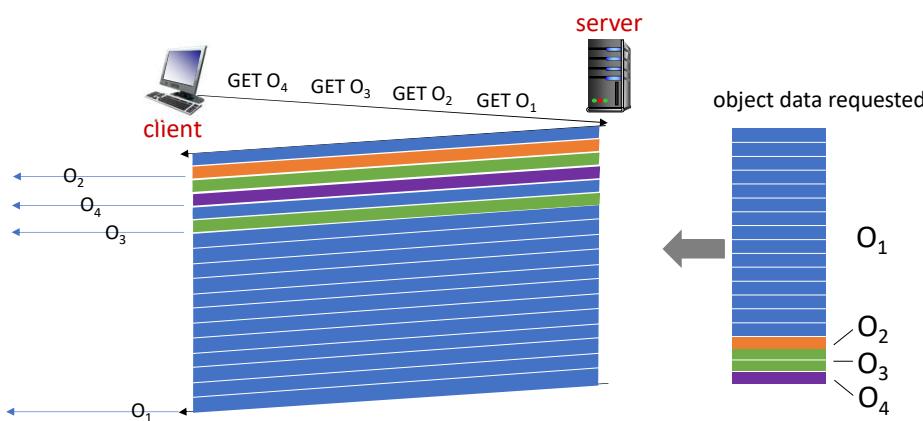


Application Layer: 2-54

54

HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved



Application Layer: 2-55

55

HTTP/2 to HTTP/3

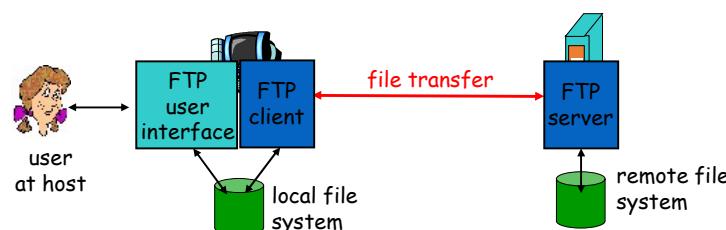
HTTP/2 over single TCP connection means:

- **recovery from packet loss** still stalls all object transmissions
 - as in HTTP 1.1, browsers have incentive to open multiple parallel TCP connections to reduce stalling, increase overall throughput
- **no security** over vanilla TCP connection
- **HTTP/3: adds security, per object error- and congestion-control (more pipelining) over UDP**
 - more on HTTP/3 in transport layer

Application Layer: 2-56

56

FTP: the File Transfer Protocol



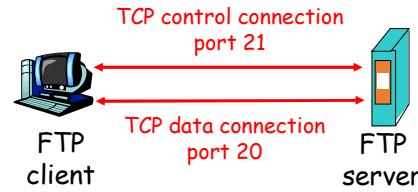
- transfer file to/from remote host
- client/server model
 - *client*: side that initiates transfer (either to/from remote)
 - *server*: remote host
- ftp: RFC 959
- ftp server: port 21

2: Application Layer 57

57

FTP: separate control, data connections

- ❑ FTP client contacts FTP server at port 21, TCP is transport protocol
- ❑ client authorized over control connection
- ❑ client browses remote directory by sending commands over control connection.
- ❑ when server receives file transfer command, server opens 2nd TCP connection (for file) to client
- ❑ after transferring one file, server closes data connection.



- ❑ server opens another TCP data connection to transfer another file.
- ❑ control connection: "out of band"
- ❑ FTP server maintains "state": current directory, earlier authentication

2: Application Layer 58

58

FTP commands, responses

Sample commands:

- sent as ASCII text over control channel
- **USER username**
- **PASS password**
- **LIST** return list of files in current directory
- **RETR filename** retrieves (gets) file
- **STOR filename** stores (puts) file onto remote host

Sample return codes

- status code and phrase (as in HTTP)
- **331 Username OK, password required**
- **125 data connection already open; transfer starting**
- **425 Can't open data connection**
- **452 Error writing file**

2: Application Layer 59

59

```

C:\> ftp
ftp> o unix.kmitl.ac.th
Connected to neutrino.crsc.kmitl.ac.th.
220-Welcome to KMITL Unix and FTP Services..
220
User <neutrino.crsc.kmitl.ac.th:<none>>: ktsakcha
331 Please specify the password.
Password:
230 Login successful.
ftp> ?
Commands may be abbreviated. Commands are:
!          delete      literal      prompt      send
?          debug       ls           put         status
append     dir        mdelete    pwd         trace
ascii      disconnect  mdirc      quit        type
bell       get        mget       quote      user
binary     glob       mkdir      recv        verbose
bye        hash       mls        remotehelp
cd         help       mput      rename
close      lcd        open       rmdir
ftp>
ftp> bye
221 Goodbye.

C:\>

```

2: Application Layer 60

60

Common FTP Commands

Command	Meaning
?	<i>to request help or information about the FTP commands</i>
ascii	<i>o set the mode of file transfer to ASCII (this is the default and transmits seven bits per character)</i>
binary	<i>to set the mode of file transfer to binary (the binary mode transmits all eight bits per byte and thus provides less chance of a transmission error and must be used to transmit files other than ASCII files)</i>
bye	<i>to exit the FTP environment (same as quit)</i>
cd	<i>to change directory on the remote machine</i>
close	<i>to terminate a connection with another computer</i>
delete	<i>to delete (remove) a file in the current remote directory (same as rm in UNIX)</i>
get	<i>to copy one file from the remote machine to the local machine</i>
help	<i>to request a list of all available FTP commands</i>
lcd	<i>to change directory on your local machine (same as UNIX cd)</i>

61

Common FTP Commands

Command	Meaning
ls	<i>to list the names of the files in the current remote directory</i>
mkdir	<i>to make a new directory within the current remote directory</i>
mget	<i>to copy multiple files from the remote machine to the local machine; you are prompted for a y/n answer before transferring each file</i>
mput	<i>to copy multiple files from the local machine to the remote machine; you are prompted for a y/n answer before transferring each file</i>
open	<i>to open a connection with another computer</i>
put	<i>to copy one file from the local machine to the remote machine</i>
pwd	<i>to find out the pathname of the current directory on the remote machine</i>
quit	<i>to exit the FTP environment (same as bye)</i>
rmdir	<i>to remove (delete) a directory in the current remote directory</i>

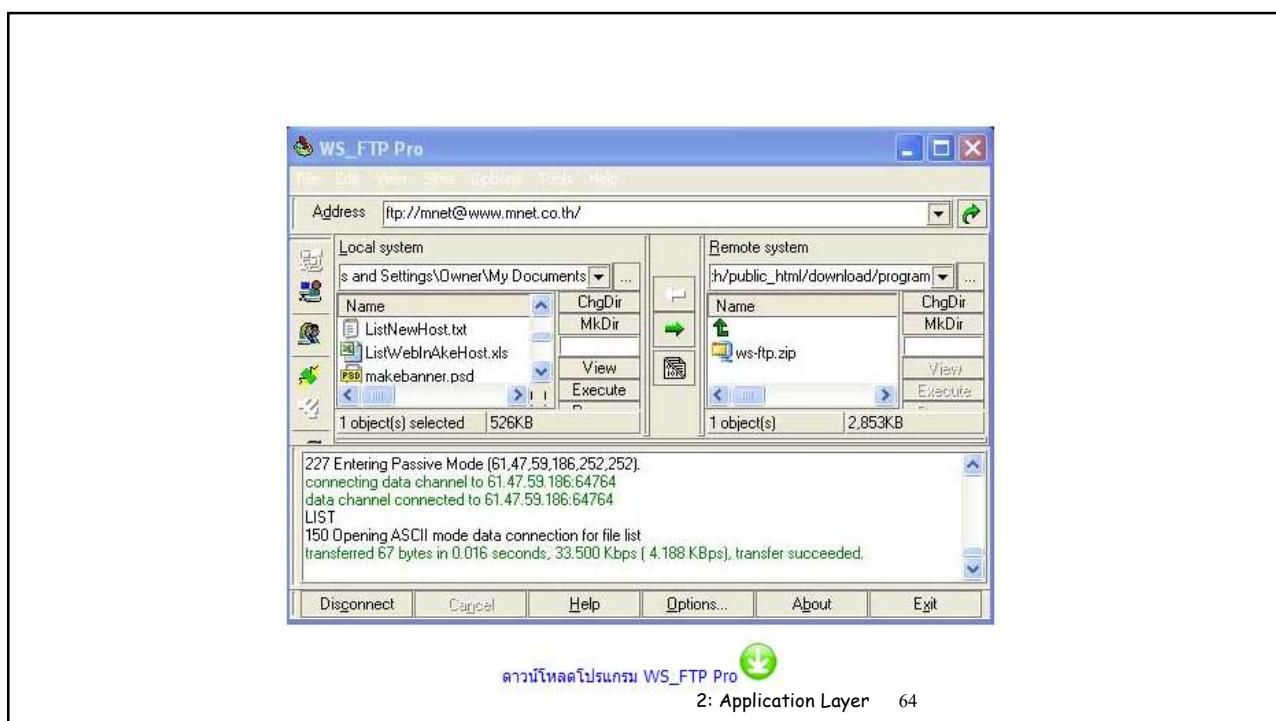
62

```

C:\WINDOWS\system32\command.com
C:\DOCUMENTS\AGNUN>ftp ftp.internic.net
Connecting to ftp.internic.net...
220 Internic FTP Server
User (ftp.internic.net:(none)): anonymous
331 Anonymous login ok, send your complete email address as your password
Password:
230 Anonymous access granted, restrictions apply
ftp> pwd
257 / is the current directory
ftp> ls -al
200 PORT command successful
150 Opening ASCII mode data connection for file list
...
domain
226 Transfer complete
ftp: 15 bytes received in 0.00Seconds 15000.00Kbytes/sec.
ftp> cd domain
250 CWD command successful
ftp> ls -al
200 PORT command successful
150 Opening ASCII mode data connection for file list
db.cache.md5
in.addr.arpa.gz
...
arpa.zone.gz.sig
ip6.arpa.gz.md5
root.zone.gz.sig
in-addr.arpa.gz.md5
.listing
named.root.md5
db.cache
db.cache.sig
edu.zone.gz.sig
named.root.md5
ip6.arpa.md5
named.cache
edu.zone.gz.md5
arpa.zone.gz
arpa.zone.gz.md5
named.cache.sig
root.zone.gz
named.root.sig
edu.zone.gz
ip6.arpa.gz
INTERNIC_ROOT_ZONE.signatures.asc
root.zone.gz.md5
named.root
INTERNIC_ROOT_ZONE.signatures
226 Transfer complete
ftp: 441 bytes received in 0.00Seconds 441000.00Kbytes/sec.
ftp> bye
221 Goodbye.
2: Application Layer 63
C:\DOCUMENTS\AGNUN>

```

63



64

Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



Application Layer: 2-66

66

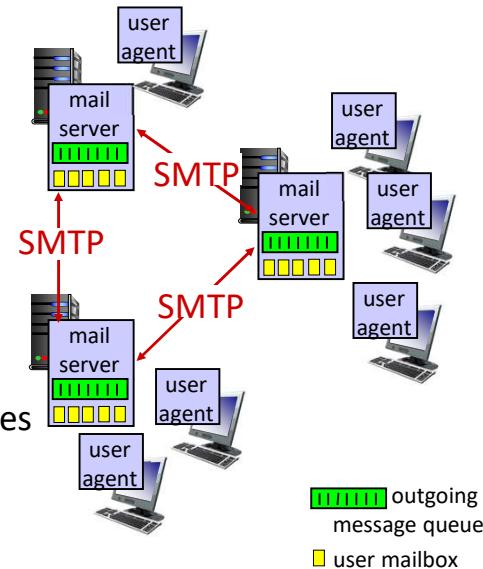
E-mail

Three major components:

- user agents
- mail servers
- simple mail transfer protocol: SMTP

User Agent

- a.k.a. “mail reader”
- composing, editing, reading mail messages
- e.g., Outlook, iPhone mail client
- outgoing, incoming messages stored on server



Application Layer: 2-67

67

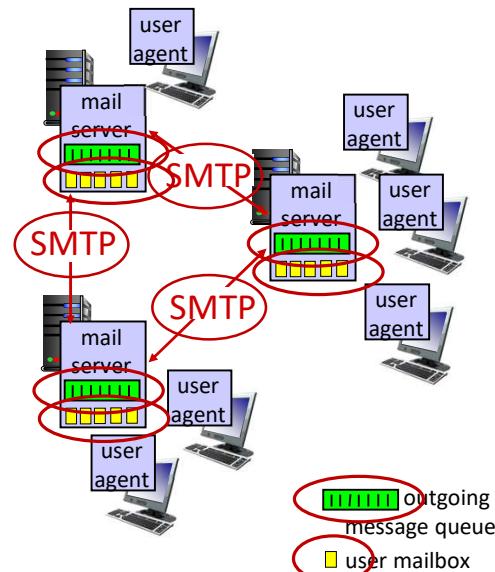
E-mail: mail servers

mail servers:

- *mailbox* contains incoming messages for user
- *message queue* of outgoing (to be sent) mail messages

SMTP protocol between mail servers to send email messages

- **client**: sending mail server
- “**server**”: receiving mail server

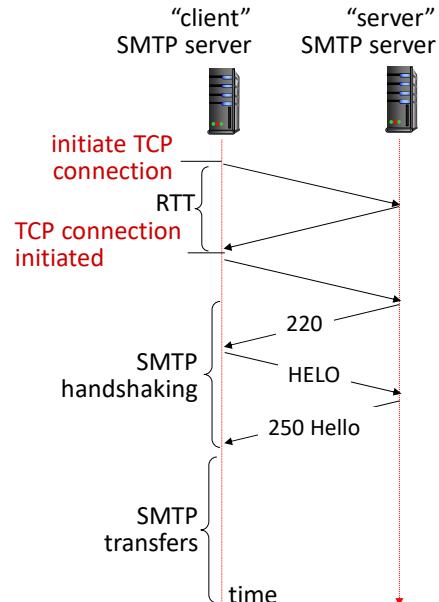


Application Layer: 2-68

68

SMTP RFC (5321)

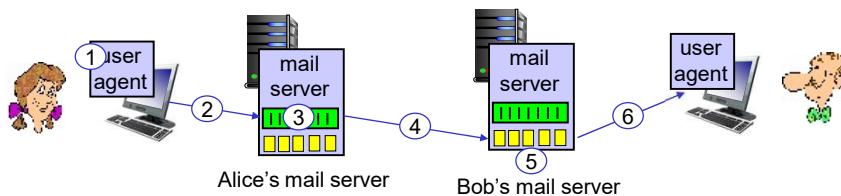
- uses **TCP** to reliably transfer email message from client (mail server initiating connection) to server, **port 25**
 - **direct transfer**: sending server (acting like client) to receiving server
- **three phases** of transfer
 - SMTP handshaking (greeting)
 - SMTP transfer of messages
 - SMTP closure
- **command/response** interaction (like HTTP)
 - commands: ASCII text
 - response: status code and phrase



69

Scenario: Alice sends e-mail to Bob

- 1) Alice uses UA to compose e-mail message "to" bob@someschool.edu
- 2) Alice's UA sends message to her mail server using SMTP; message placed in message queue
- 3) client side of SMTP at mail server opens TCP connection with Bob's mail server
- 4) SMTP client sends Alice's message over the TCP connection
- 5) Bob's mail server places the message in Bob's mailbox
- 6) Bob invokes his user agent to read message



Application Layer: 2-70

70

Sample SMTP interaction

S: 220 **hamburger.edu**

Application Layer: 2-71

71

Try SMTP interaction for yourself:

- **telnet servername 25**
- see 220 reply from server
- enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands
above lets you send email without using email client (reader)

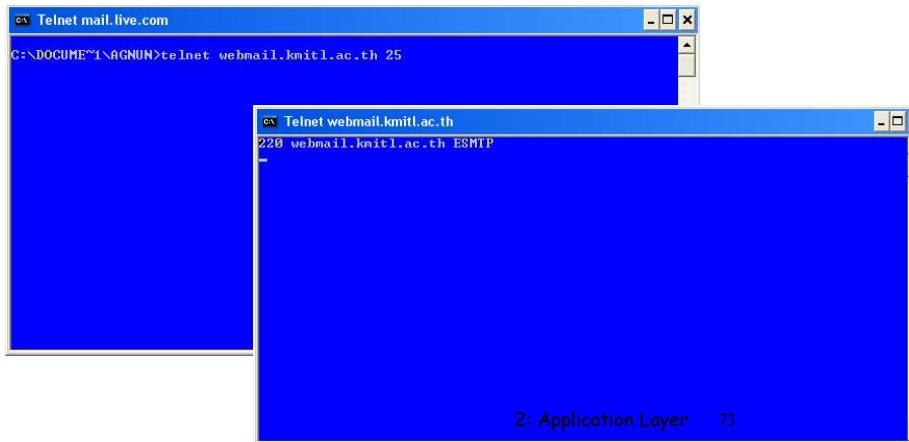
2: Application Layer 72

72

33

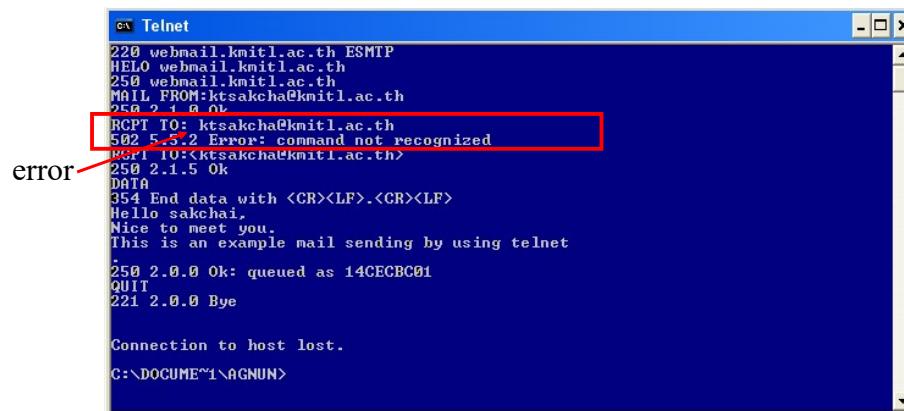
An Example SMTP Session

- o How to connect to an SMTP server?
 - o telnet servername 25
 - o A TCP connection gets established over port number 25
 - o The telnet client and the mail server can start a dialogue



73

An Example SMTP Session (cont.)



2: Application Layer 74

74

An Example SMTP Session (cont.)

```

220 webmail.kmitl.ac.th ESMTP
HELO webmail.kmitl.ac.th
250 webmail.kmitl.ac.th
MAIL FROM:ktsakcha@kmitl.ac.th
250 2.1.0 Ok
RCPT TO:ktsakcha@kmitl.ac.th] Recipient 1
250 2.1.5 Ok
RCPT TO:<kusomsak@kmitl.ac.th> Recipient 2
250 2.1.5 Ok
RCPT TO:<kvaranya@kmitl.ac.th> Recipient 3
250 2.1.5 Ok
DATA
354 End data with <CR><LF>.<CR><LF>
Hello,
This is an example mail sending by using telnet program
250 2.0.0 Ok: queued as 55EFEB01
QUIT
221 2.0.0 Bye

Connection to host lost.
C:\DOCUMENTS\AGNUNO>

```

Sending email to several recipients

2: Application Layer 75

75

SMTP: observations

comparison with HTTP:

- **HTTP:** client **pull**
- **SMTP:** client **push**
- both have **ASCII command/response** interaction, **status codes**
- **HTTP:** **each object** encapsulated in its own response message
- **SMTP:** multiple objects sent in **multipart message**
- **SMTP uses persistent connections**
- **SMTP requires message (header & body) to be in 7-bit ASCII**
- **SMTP server uses CRLF,CRLF to determine end of message**

Application Layer: 2-76

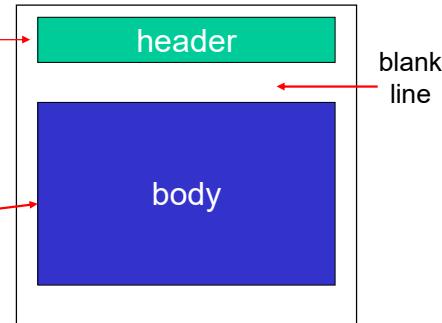
76

Mail message format

SMTP: protocol for exchanging e-mail messages, defined in [RFC 5321](#)
(like [RFC 7231](#) defines [HTTP](#))

RFC 2822 defines *syntax* for e-mail message itself (like HTML defines syntax for web documents)

- header lines, e.g., _____
 - To:
 - From:
 - Subject:these lines, within the body of the email message area different from ~~SMTP MAIL FROM~~ ~~RCPT TO~~: commands!
 - Body: the “message”, ASCII characters only



Application Layer: 2-77

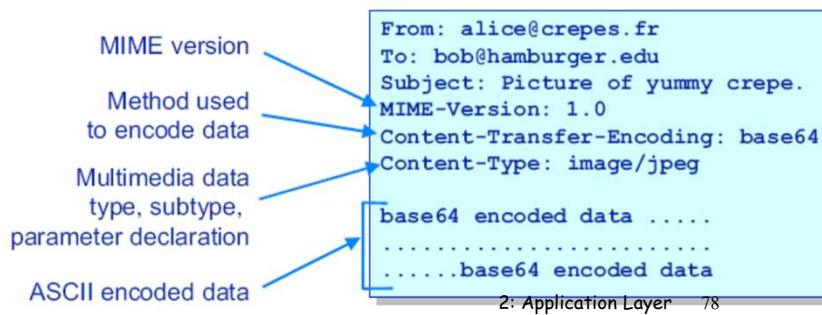
77

Mail Message Format

MIME – Multipurpose Internet Mail Extensions

(RFC 2045, 2056)

- Multimedia mail extensions
 - SMTP requires all data to be *7-bit ASCII characters*
 - All non-ASCII data must be encoded as ASCII strings
 - Additional lines in the message header declare MIME content type



2: Application Layer 78

78

MIME types and subtypes

Type	Subtype	Description
Text	Plain html	Unformatted Text Html format
Image	Gif Jpeg bmp	Still picture in GIF format Still picture in JPEG format Still picture in BMP format
Audio	Basic X-wav	Audible sound (au) Wave file
Video	Mpeg Quicktime	Movie in MPEG format Mov.
Multipart	Mixed	Independent parts in the specified order

80

80

MIME Types : *Multipart Types*

MIME version

```

From: alice@crepes.fr
To: bob@hamburger.edu
Subject: Picture of yummy crepe.
MIME-Version: 1.0
Content-Type: multipart/mixed; boundary=98766789

```

Header/Body separators

```

--98766789
Content-Transfer-Encoding: quoted-printable
Content-Type: text/plain

Dear Bob,
Please find a picture of a crepe.
--98766789
Content-Transfer-Encoding: base64
Content-Type: image/jpeg

base64 encoded data .....
.....
.....base64 encoded data
--98766789--

```

2: Application Layer 82

82

Received Message

- Receiving Server ที่ได้รับ message ด้วย RFC 822 และ MIME header line จะทำการเพิ่ม
 - Received : header line ในบรรทัดบนสุดของ message

Received: from crepes.fr by hamburger.edu; 12 Oct 98 15:27:39 GMT

From: alice@crepes.fr
 To: bob@hamburger.edu
 Subject: Picture of yummy crepe.
 MIME-Version: 1.0
 Content-Transfer-Encoding: base64
 Content-Type: image/jpeg

2: Application Layer 83

83

Received Message

- บางครั้ง บรรทัด Received: มีหลายบรรทัด เช่น

Received: from hamburger.edu by sushi.jp; 3 Jul 01 15:30:01 GMT
Received: from crepes.fr by hamburger.edu; 3 Jul 01 15:17:39 GMT

- เกิดจาก user มีการ forward mail ไปยัง SMTP Server อื่นๆ

2: Application Layer 84

84

Base64 Encoding Table

Base64 Encoding Table

Value	Binary	Char									
0	000000	A	16	010000	Q	32	100000	g	48	110000	w
1	000001	B	17	010001	R	33	100001	h	49	110001	x
2	000010	C	18	010010	S	34	100010	i	50	110010	y
3	000011	D	19	010011	T	35	100011	j	51	110011	z
4	000100	E	20	010100	U	36	100100	k	52	110100	0
5	000101	F	21	010101	V	37	100101	l	53	110101	1
6	000110	G	22	010110	W	38	100110	m	54	110110	2
7	000111	H	23	010111	X	39	100111	n	55	110111	3
8	001000	I	24	011000	Y	40	101000	o	56	111000	4
9	001001	J	25	011001	Z	41	101001	p	57	111001	5
0	001010	K	26	011010	a	42	101010	q	58	111010	6
11	001011	L	27	011011	b	43	101011	r	59	111011	7
12	001100	M	28	011100	c	44	101100	s	60	111100	8
13	001101	N	29	011101	d	45	101101	t	61	111101	9
14	001110	O	30	011110	e	46	101110	u	62	111110	+
15	001111	P	31	011111	f	47	101111	v	63	111111	/

2: Application Layer

85

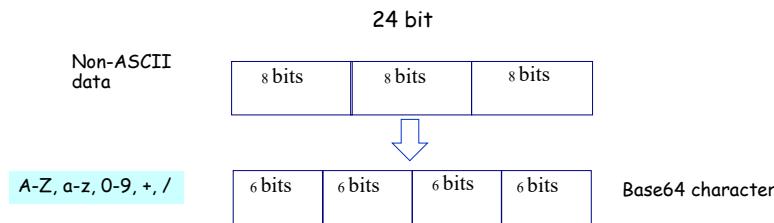
85

Base64 Encoding

- ทำการเข้ารหัสครั้งละ 3 ตัวอักษร โดยทำการแปลงแต่ละตัวอักษรเป็นเลขฐานสองขนาด 8 บิต โดยเทียบจากตาราง ASCII (ให้เติม 0 ตรงตำแหน่งบิตซ้ายสุด ให้เป็น 8 บิต)
- นำเลขฐานสองของแต่ละตัวอักษรมาเขียนเรียงกัน (24 บิต)
- ทำการแบ่งออกเป็นชุดๆ ละ 6 บิต (ได้ทั้งสิ้น 4 ชุด)
- ทำการแปลงแต่ละชุดให้เป็นตัวอักษรโดยเทียบจากตาราง Base64
- กลับไปทำขั้นตอนที่ 1 ใหม่ จนกระทั่ง เหลือข้อมูลกลุ่มสุดท้าย ซึ่งมีໄได้ 3 กรณีดังนี้
 - กรณีที่ 1 ข้อมูลมี 3 ตัวอักษร致命 (ทำขั้นตอนที่ 1 ลีบ 4)
 - กรณีที่ 2 ข้อมูลมี จำนวน 2 ตัวอักษร致命
 - กรณีที่ 3 ข้อมูลมี จำนวน 1 ตัวอักษร致命
- ทำการแปลงตัวอักษรที่เหลือ (2 ตัวหรือ 1 ตัว) ให้เป็นเลขฐานสอง
- เติม 0 ให้ครบ 24 บิต
- ทำการแบ่งออกเป็นชุดๆ ละ 6 บิต (ได้ทั้งสิ้น 4 ชุด)
- ทำการแปลงแต่ละชุดให้เป็นตัวอักษร โดยเทียบจากตาราง Base64 ถ้าชุดใดเป็น 0 ทั้งหมดให้แทนด้วย

88

Base64 Encoding



ข้อมูลกู้มุดหัก ชั่งมีได้ 3 กรณีดังนี้

- กรณีที่ 1 ข้อมูลมี 3 ไบต์พอดี (ทำขั้นตอนที่ 1 ถึง 3)
- กรณีที่ 2 ข้อมูลมี จำนวน 2 ไบต์ (ขาด 1 ไบต์ ครบ 3 ไบต์)
- กรณีที่ 3 ข้อมูลมี จำนวน 1 ไบต์ (ขาด 2 ไบต์ ครบ 3 ไบต์)

- ❖ กรณีที่ 2 และ 3 ต้อง 0 ไบต์ครบ 24 บิต
- ❖ ทำการแบ่งออกเป็นชุดๆ ละ 6 บิต (ได้ทั้งสิ้น 4 ชุด)
- ❖ ทำการแปลงค่าของแต่ละชุดให้เป็นตัวอักษร โดยพิมพ์จากตาราง Base64 ดังภาพไปใน 2: Application Layer ทั้งหมดให้แนบตัวอย่าง “=” (กรณีที่ 2 จะมี “=” กรณีที่ 3 จะมี “==”)

90

กรณีที่ 1 ครบ 3 ไบต์ หรือ 24 บิต

○ จงแปลง ABC ให้อยู่ในรูปของ Base64

○ $A=41_{16}$, $B=42_{16}$, $C=43_{16}$

○ 0100 0001 0100 0010 0100 0011

○ 0100 0001 0100 0010 0100 0011

○ Q U J D

○ $ABC \rightarrow QUJD$

กรณี 2 เหลือ 2 ไบต์

- จงแปลง ABCDE ให้อยู่ในรูปของ Base64
- $A=41_{16}$, $B=42_{16}$, $C=43_{16}$, $D=44_{16}$, $E=45_{16}$
- 0100 0001 0100 0010 0100 0011 01000100 01000101
- 0100 0001 0100 0010 0100 0011 01000100 01000101 00000000
- Q U J D R E U =
- ABCDE \rightarrow QUJDREU=

2: Application Layer 94

94

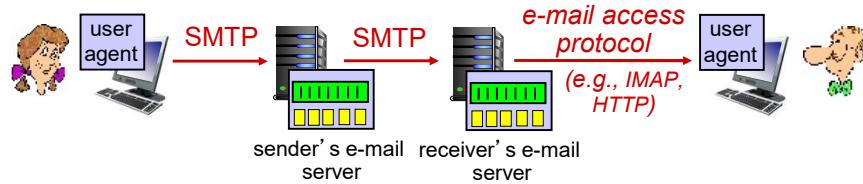
กรณี 3 เหลือ 1 ไบต์

- จงแปลง ABCD ให้อยู่ในรูปของ Base64
- $A=41_{16}$, $B=42_{16}$, $C=43_{16}$, $D=44_{16}$
- 0100 0001 0100 0010 0100 0011 01000100
- 0100 0001 0100 0010 0100 0011 01000100 0000000000000000
- Q U J D R A = =
- ABCD \rightarrow QUJDRA==

2: Application Layer 95

95

Retrieving email: mail access protocols

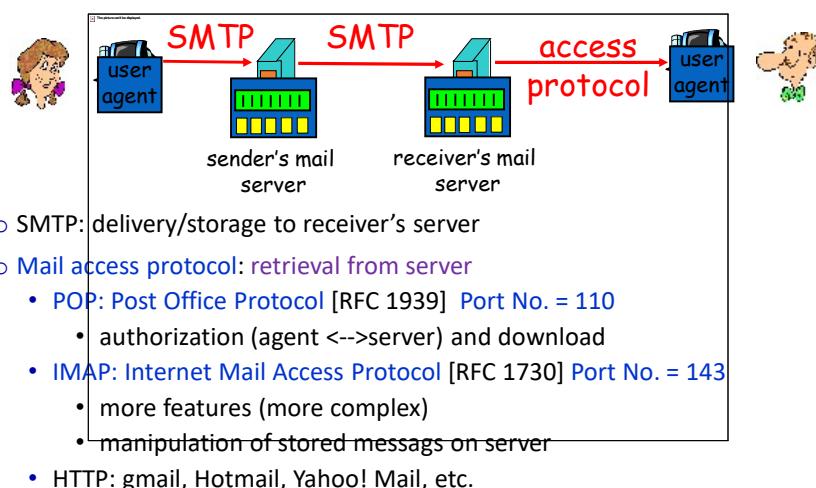


- **SMTP:** delivery/storage of e-mail messages to receiver's server
- mail access protocol: retrieval from server
 - **IMAP:** Internet Mail Access Protocol [RFC 3501]: messages stored on server, IMAP provides retrieval, deletion, folders of stored messages on server
- **HTTP:** gmail, Hotmail, Yahoo!Mail, etc. provides web-based interface on top of STMP (to send), IMAP (or POP) to retrieve e-mail messages

Application Layer: 2-96

96

Mail access protocols



2: Application Layer 97

97

POP3 protocol

authorization phase

- o client commands:
 - **user**: declare username
 - **pass**: password
- o server responses
 - **+OK**
 - **-ERR**

transaction phase, client:

- o **list**: list message numbers
- o **retr**: retrieve message by number
- o **delete**: delete
- o **quit**

```

S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off

```

2: Application Layer 98

98

POP3 (more) and IMAP

More about POP3

- o Previous example uses “download and delete” mode.
- o Bob **cannot** re-read e-mail if he changes client
- o “Download-and-keep”: copies of messages on different clients
- o POP3 is stateless across sessions

IMAP

- o Keep all messages in one place: the **server**
- o Allows user to organize messages in **folders**
- o IMAP keeps user state across sessions:
 - o **names of folders** and mappings between message IDs and folder name

2: Application Layer 99

99

Application Layer: Overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



Application Layer: 2-100

100

DNS: Domain Name System

people: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., cs.umass.edu - used by humans

Q: how to map between IP address and name, and vice versa ?

Domain Name System (DNS):

- *distributed database* implemented in hierarchy of many *name servers*
- *application-layer protocol:* hosts, DNS servers communicate to *resolve* names (address/name translation)
 - *note:* core Internet function, **implemented as application-layer protocol**
 - complexity at network’s “edge”

Application Layer: 2-101

101

DNS: services, structure

DNS services:

- hostname-to-IP-address translation
- host aliasing
 - canonical, alias names
- mail server aliasing
- load distribution
 - replicated Web servers: many IP addresses correspond to one name

Q: Why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

A: doesn't scale!

- Comcast DNS servers alone: 600B DNS queries/day
- Akamai DNS servers alone: 2.2T DNS queries/day

Application Layer: 2-102

102

Thinking about the DNS

humongous distributed database:

- ~ billion records, each simple

handles many *trillions* of queries/day:

- *many* more reads than writes
- *performance matters*: almost every Internet transaction interacts with DNS - msecs count!

organizationally, physically decentralized:

- millions of different organizations responsible for their records

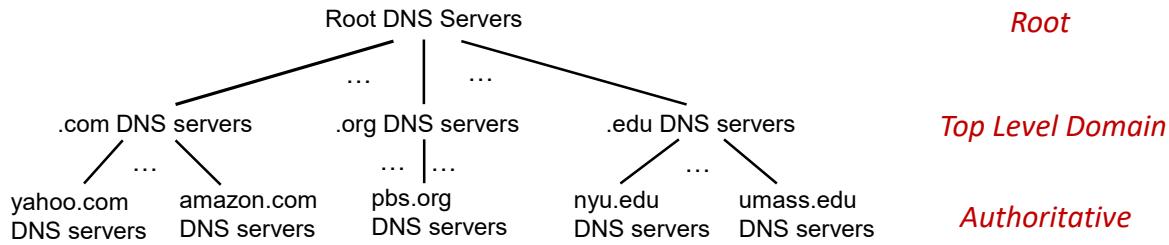
“bulletproof”: reliability, security



Application Layer: 2-103

103

DNS: a distributed, hierarchical database



Client wants IP address for www.amazon.com; 1st approximation:

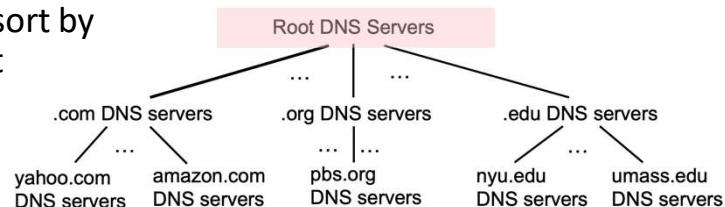
- client queries **root server** to find **.com** DNS server
- client queries **.com** DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get **IP address** for www.amazon.com

Application Layer: 2-104

104

DNS: root name servers

- official, contact-of-last-resort by name servers that can not resolve name



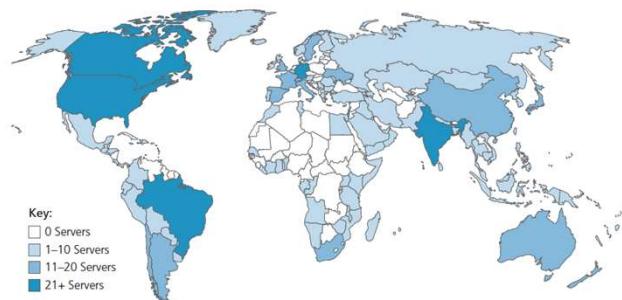
Application Layer: 2-105

105

DNS: root name servers

- official, contact-of-last-resort by name servers that can not resolve name
- *incredibly important* Internet function
 - Internet couldn't function without it!
 - DNSSEC – provides security (authentication, message integrity)
- ICANN (Internet Corporation for Assigned Names and Numbers) manages root DNS domain

13 logical root name “servers” worldwide each “server” replicated many times (~200 servers in US)



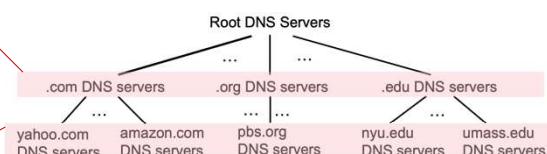
Application Layer: 2-106

106

Top-Level Domain, and authoritative servers

Top-Level Domain (TLD) servers:

- responsible for .com, .org, .net, .edu, .aero, .jobs, .museums, and all top-level country domains, e.g.: .cn, .uk, .fr, .ca, .jp
- Network Solutions: authoritative registry for .com, .net TLD
- Educause: .edu TLD



authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

Application Layer: 2-107

107

Local DNS name servers

- when **host** makes **DNS query**, it is sent to its **local DNS server**
 - Local DNS server returns reply, answering:
 - from its **local cache** of recent name-to-address translation pairs (possibly **out of date!**)
 - **forwarding request** into **DNS hierarchy** for resolution
 - each ISP has local DNS name server; to find yours:
 - MacOS: % scutil --dns
 - Windows: >ipconfig /all
- **local DNS server doesn't** strictly **belong to hierarchy**

Application Layer: 2-108

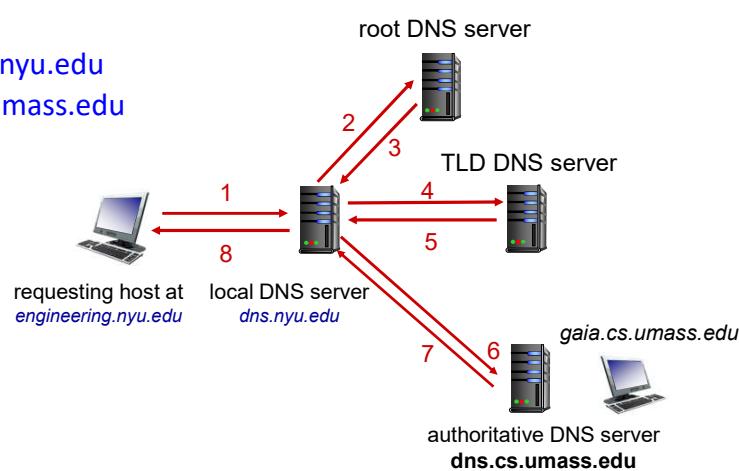
108

DNS name resolution: iterated query

Example: host at engineering.nyu.edu wants IP address for gaia.cs.umass.edu

Iterated query:

- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



Application Layer: 2-109

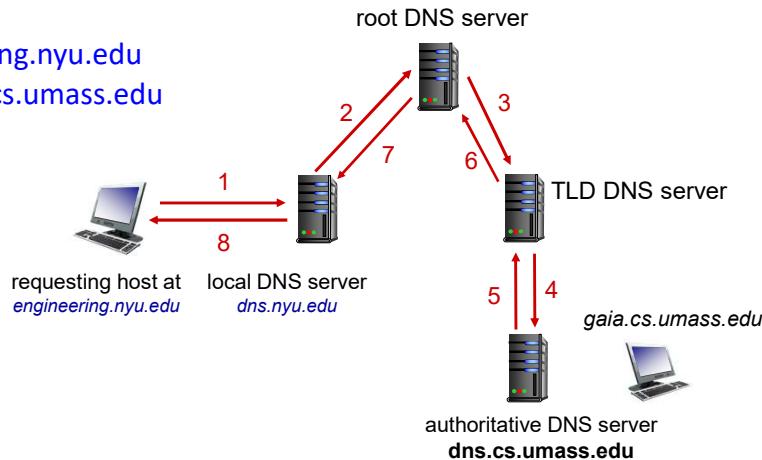
109

DNS name resolution: recursive query

Example: host at engineering.nyu.edu
wants IP address for gaia.cs.umass.edu

Recursive query:

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?



Application Layer: 2-110

110

Caching DNS Information

- once (any) name server learns mapping, it *caches* mapping, and *immediately* returns a cached mapping in response to a query
 - caching improves *response time*
 - cache entries *timeout* (disappear) after some *time (TTL)*
 - TLD servers typically cached in *local name servers*
- cached entries may be *out-of-date*
 - if named host *changes IP address*, may *not be known* Internet-wide until all TTLs expire!
 - *best-effort name-to-address translation!*

Application Layer: 2-111

111

DNS records

TTL : time to live of RR;
It determines when a resource should be removed from cache

DNS: distributed database storing resource records (RR)

RR format: (name, value, type, ttl)

type=A

- name is hostname
- value is IP address

type=NS

- name is domain (e.g., foo.com)
- value is hostname of authoritative name server for this domain

type=CNAME

- name is alias name for some “canonical” (the real) name
- www.ibm.com is really servereast.backup2.ibm.com
- value is canonical name

type=MX

- value is canonical name of SMTP mail server associated with name

Application Layer: 2-112

112

DNS protocol messages

DNS *query* and *reply* messages, both have same *format*:

message header:

- identification: 16 bit # for query, reply to query uses same #
- flags:
 - query or reply
 - recursion desired
 - recursion available
 - reply is authoritative

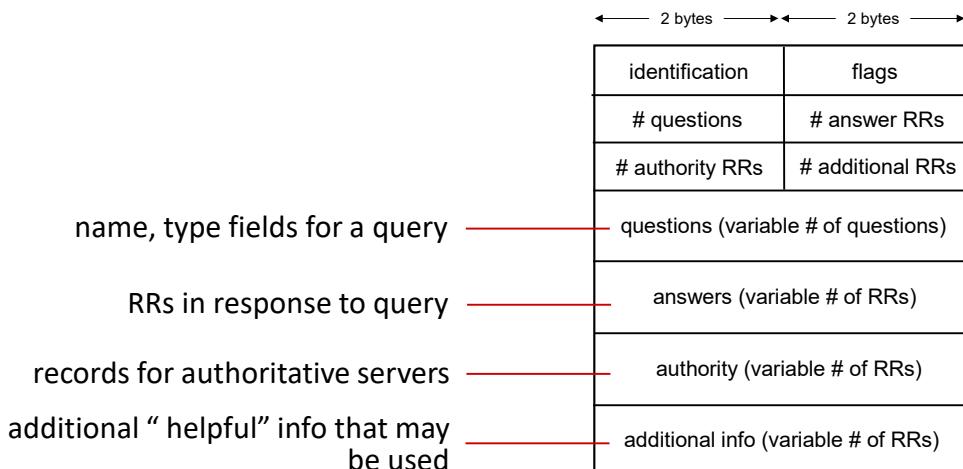
identification	flags
# questions	# answer RRs
# authority RRs	# additional RRs
questions (variable # of questions)	
answers (variable # of RRs)	
authority (variable # of RRs)	
additional info (variable # of RRs)	

Application Layer: 2-113

113

DNS protocol messages

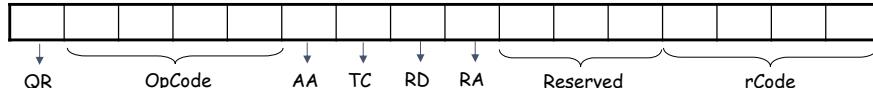
DNS *query* and *reply* messages, both have same *format*:



Application Layer: 2-114

114

Flags Field



- o QR (Query/Response) defines type of message
 - o If it is 0, the message is a **query**
 - o If it is 1, the message is a **response**
- o OpCode : defines type of query or response
 - 0 if standard query, 1 if inverse query, 2 if a server status request
- o AA (authoritative answer)
 - o When it is set (value of 1) it means that name server is an authoritative server
 - o It is used only in **response** message
- o TC (truncated)
 - o When it is set (value of 1), it means that **response** was more than 512 bytes and truncated to 512
 - o It is used when DNS uses services of **UDP**
- o RD (recursion desired)
 - o When it is **set** (value of 1), it means client desires **recursive** answer
 - o It is **set in query message and repeated in response message**
- o RA (recursion available), this bit is valid in **response** and denotes whether recursive query support is available
- o Reserved
 - o 3-bit subfield is set to 000
- o rCode shows status of error in **response**

2: Application Layer 115

115

Flags Field

Value	Meaning
0	No error
1	Format error
2	Problem at name server
3	Domain reference problem
4	Query type not supported
5	Administratively prohibited
6-15	Reserved

Value of rCode

2: Application Layer 116

116

Getting your info into the DNS

example: new startup “[Network Utopia](#)”

- register name [networkuptopia.com](#) at *DNS registrar* (e.g., Network Solutions)
 - provide names, IP addresses of authoritative name server (primary and secondary)
 - registrar inserts NS, A RRs into .com TLD server:
 (networkutopia.com, dns1.networkutopia.com, NS)
 (dns1.networkutopia.com, 212.212.212.1, A)
- create authoritative server locally with IP address 212.212.212.1
 - type A record for www.networkuptopia.com
 - type MX record for networkutopia.com

Application Layer: 2-117

117

Encapsulation

- o DNS can use either **UDP** or **TCP**.
- o In both cases the well-known port used by the server is **port 53**.
- o **UDP** is used when the size of **response message** is **less than 512 bytes** because most UDP packages have a 512-byte packet size limit.
- o If the size of **response message** is **more than 512 bytes**, a **TCP connection** is used.

2: Application Layer 118

118

DNS security

DDoS attacks

- bombard root servers with traffic
 - not successful to date
 - traffic filtering
 - **local DNS servers** cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers
 - potentially more dangerous

Spoofing attacks

- intercept DNS queries, returning bogus replies
 - DNS cache poisoning
 - RFC 4033: DNSSEC authentication services

Application Layer: 2-119

119

Application Layer: Overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP

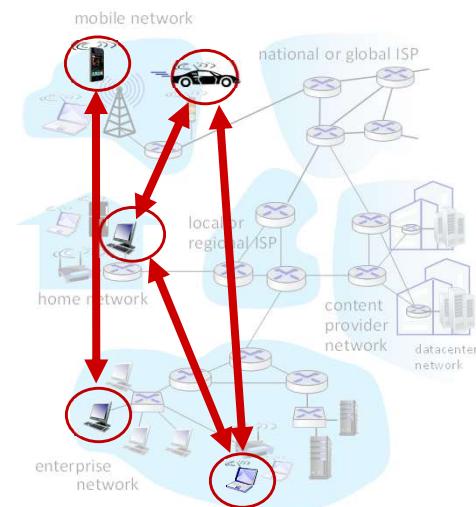


Application Layer: 2-120

120

Peer-to-peer (P2P) architecture

- *no always-on server*
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, and new service demands
- peers are intermittently connected and change IP addresses
 - complex management
- examples: P2P file sharing (BitTorrent), streaming (KanKan), VoIP (Skype)



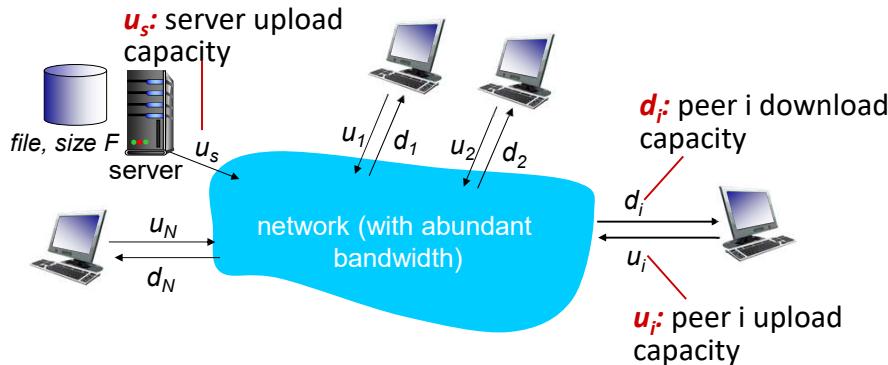
Application Layer: 2-121

121

File distribution: client-server vs P2P

Q: how much time to distribute file (size F) from one server to N peers?

- peer upload/download capacity is limited resource



Introduction: 1-122

122

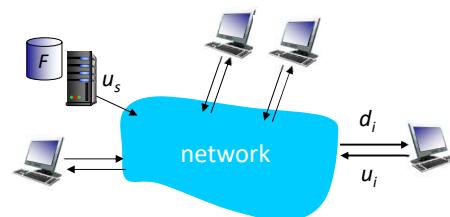
File distribution time: client-server

- **server transmission:** must sequentially send (upload) N file copies:

- time to send one copy: F/u_s
- time to send N copies: NF/u_s

- **client:** each client must download file copy

- d_{min} = min client download rate
- min client download time: F/d_{min}



$$\text{time to distribute } F \text{ to } N \text{ clients using client-server approach} \quad D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

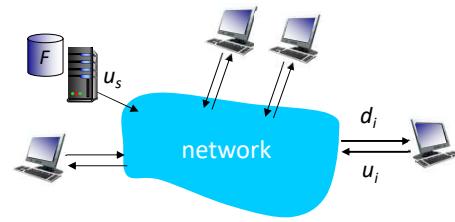
increases linearly in N

Introduction: 1-123

123

File distribution time: P2P

- **server transmission:** must upload at least one copy:
 - time to send one copy: F/u_s
- **client:** each client must download file copy
 - min client download time: F/d_{min}
- **clients:** as aggregate must download NF bits
 - max upload rate (limiting max download rate) is $u_s + \sum u_i$



*time to distribute F
to N clients using
P2P approach*

$$D_{P2P} \geq \max\{F/u_s, F/d_{min}, NF/(u_s + \sum u_i)\}$$

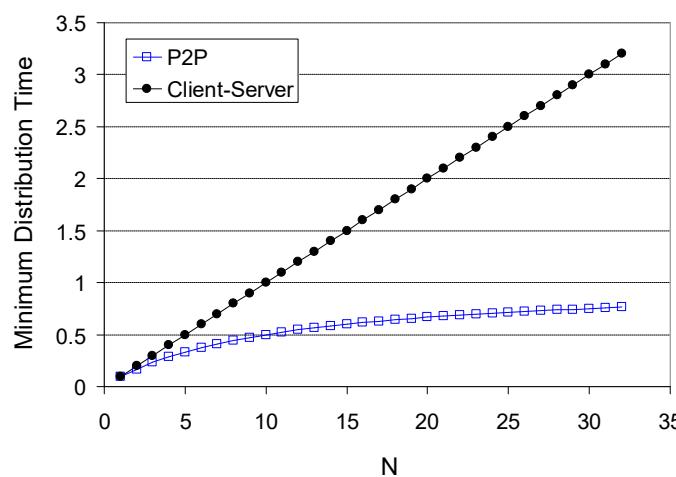
increases linearly in N ...
... but so does this, as each peer brings service capacity

Application Layer: 2-124

124

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$

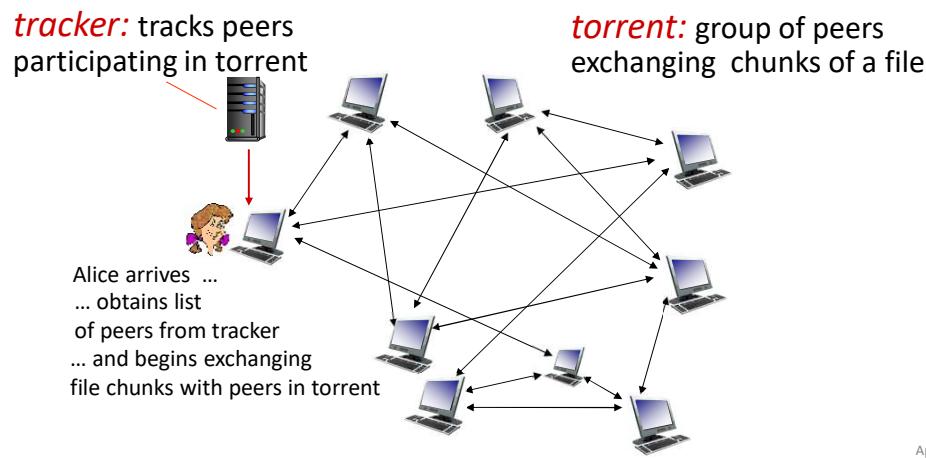


Application Layer: 2-125

125

P2P file distribution: BitTorrent

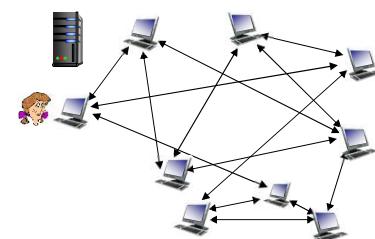
- file divided into 256Kb chunks
- peers in torrent send/receive file chunks



126

P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn:** peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



Application Layer: 2-127

127

BitTorrent: requesting, sending file chunks

Requesting chunks:

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

Sending chunks: tit-for-tat

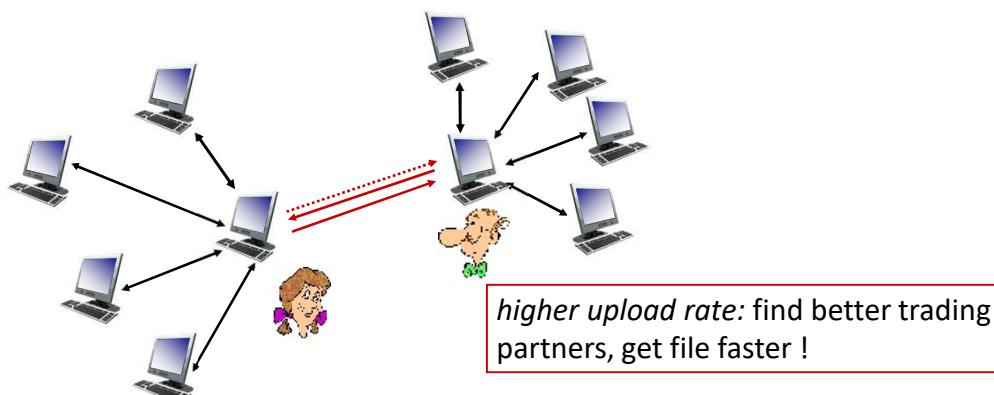
- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4

Application Layer: 2-128

128

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Application Layer: 2-129

129

Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



Application Layer: 2-130

130

Video Streaming and CDNs: context

- stream video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube, Amazon Prime: 80% of residential ISP traffic (2020)
- *challenge:* scale - how to reach ~1B users?
- *challenge:* heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution:* distributed, application-level infrastructure



Application Layer: 2-131

131

Multimedia: video

- video: sequence of images displayed at constant rate
 - e.g., 24 images/sec
- digital image: array of pixels
 - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i



frame $i+1$

Application Layer: 2-132

132

Multimedia: video

- **CBR: (constant bit rate):** video encoding rate fixed
- **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes
- **examples:**
 - MPEG 1 (CD-ROM) 1.5 Mbps
 - MPEG2 (DVD) 3-6 Mbps
 - MPEG4 (often used in Internet, 64Kbps – 12 Mbps)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i



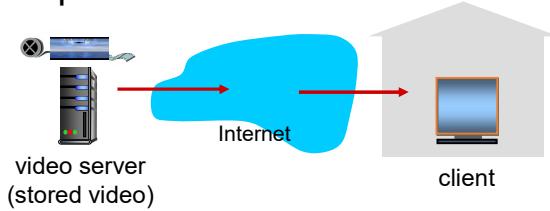
frame $i+1$

Application Layer: 2-133

133

Streaming stored video

simple scenario:



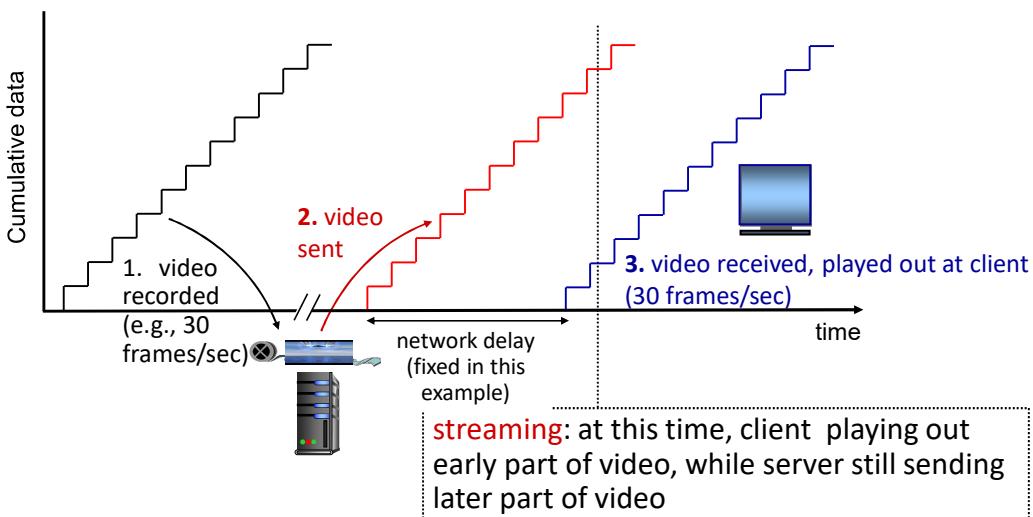
Main challenges:

- server-to-client bandwidth will **vary** over time, with changing network congestion levels (in house, access network, network core, video server)
- packet loss, delay due to congestion will delay playout, or result in poor video quality

Application Layer: 2-134

134

Streaming stored video



Application Layer: 2-135

135

Streaming stored video: challenges

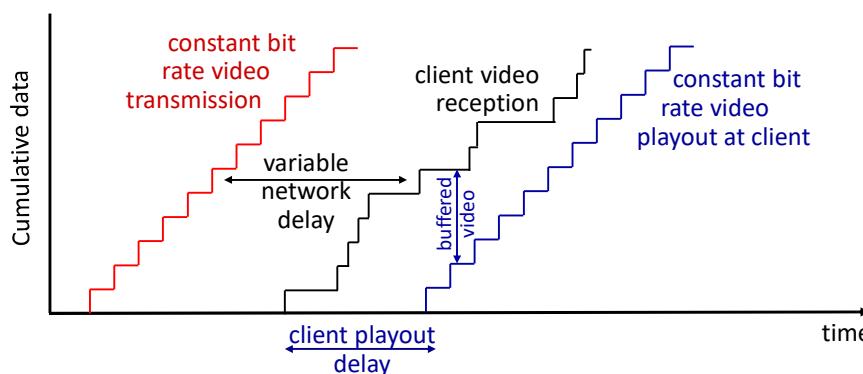
- **continuous playout constraint:** during client video playout, playout timing must match original timing
 - ... but **network delays are variable** (jitter), so will need **client-side buffer** to match continuous playout constraint
- **other challenges:**
 - client interactivity: pause, fast-forward, rewind, jump through video
 - video packets may be lost, retransmitted



Application Layer: 2-136

136

Streaming stored video: playout buffering



- **client-side buffering and playout delay:** compensate for network-added delay, delay jitter

Application Layer: 2-137

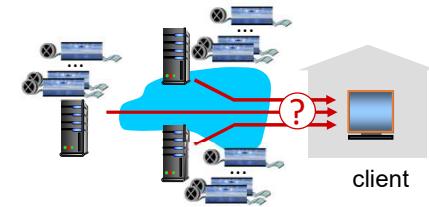
137

Streaming multimedia: DASH

Dynamic, Adaptive Streaming over HTTP

server:

- divides video file into multiple chunks
- each chunk encoded at multiple different rates
- different rate encodings stored in different files
- files replicated in various CDN nodes
- *manifest file*: provides URLs for different chunks



client:

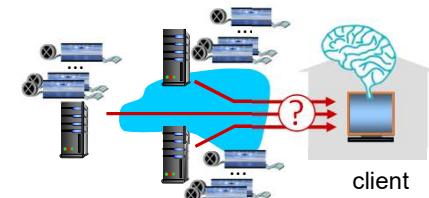
- periodically estimates server-to-client bandwidth
- consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time), and from different servers

Application Layer: 2-138

138

Streaming multimedia: DASH

- “*intelligence*” at client: client determines
 - *when* to request chunk (so that buffer starvation, or overflow does not occur)
 - *what encoding rate* to request (higher quality when more bandwidth available)
 - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)



Streaming video = encoding + DASH + playout buffering

Application Layer: 2-139

139

Content distribution networks (CDNs)

challenge: how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

- **option 1:** single, large “mega-server”
 - single point of failure
 - point of network congestion
 - long (and possibly congested) path to distant clients

....quite simply: this solution *doesn't scale*

Application Layer: 2-140

140

Content distribution networks (CDNs)

challenge: how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

- **option 2:** store/serve multiple copies of videos at multiple geographically distributed sites (*CDN*)
 - *enter deep:* push CDN servers deep into many access networks
 - close to users
 - Akamai: 240,000 servers deployed in > 120 countries (2015)
 - *bring home:* smaller number (10's) of larger clusters in POPs near access nets
 - used by Limelight

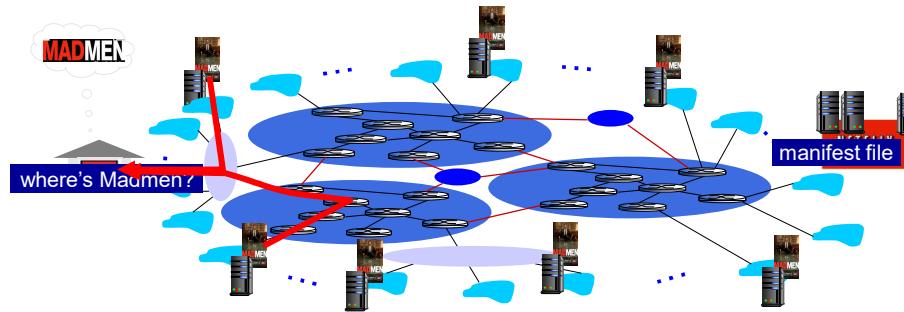


Application Layer: 2-141

141

Content distribution networks (CDNs)

- CDN: stores copies of content (e.g. MADMEN) at CDN nodes
- subscriber requests content, service provider returns manifest
 - using manifest, client retrieves content at highest supportable rate
 - may choose different rate or copy if network path congested



Application Layer: 2-142

142

Content distribution networks (CDNs)



OTT challenges: coping with a congested Internet from the “edge”

- what content to place in which CDN node?
- from which CDN node to retrieve content? At which rate?

Application Layer: 2-143

143

Application Layer: Overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



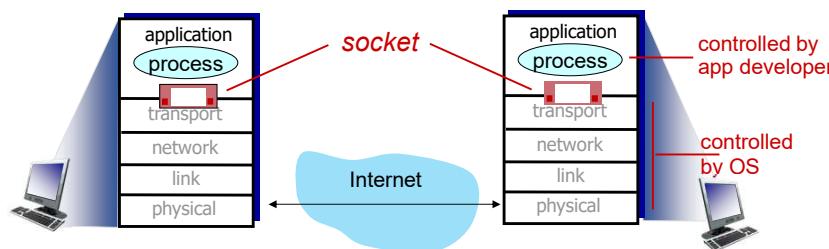
Application Layer: 2-144

144

Socket programming

goal: learn how to build client/server applications that communicate using sockets

socket: door between application process and end-end-transport protocol



Application Layer: 2-145

145

Socket programming

Two socket types for two transport services:

- **UDP:** unreliable datagram
- **TCP:** reliable, byte stream-oriented

Application Example:

1. client reads a line of characters (data) from its keyboard and sends data to server
2. server receives the data and converts characters to uppercase
3. server sends modified data to client
4. client receives modified data and displays line on its screen

Application Layer: 2-146

146

Socket programming with UDP

UDP: no “connection” between client and server:

- no handshaking before sending data
- sender explicitly attaches IP destination address and port # to each packet
- receiver extracts sender IP address and port# from received packet

UDP: transmitted data may be lost or received out-of-order

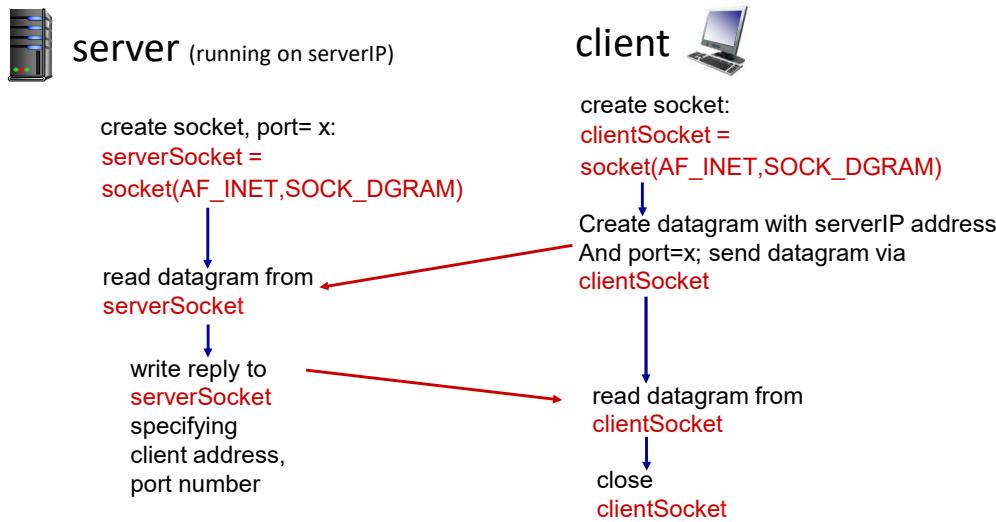
Application viewpoint:

- UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server processes

Application Layer: 2-147

147

Client/server socket interaction: UDP



Application Layer: 2-148

148

Example app: UDP client

Python UDPClient

```

include Python's socket library → from socket import *
serverName = 'hostname'
serverPort = 12000
create UDP socket for server → clientSocket = socket(AF_INET,
                                                       SOCK_DGRAM)
get user keyboard input → message = raw_input('Input lowercase sentence:')
attach server name, port to message; send into socket → clientSocket.sendto(message.encode(),
                           (serverName, serverPort))
read reply characters from socket into string → modifiedMessage, serverAddress =
                                               clientSocket.recvfrom(2048)
print out received string and close socket → print modifiedMessage.decode()
                                              clientSocket.close()

```

Application Layer: 2-149

149

Example app: UDP server

Python UDPServer

```

from socket import *
serverPort = 12000
create UDP socket → serverSocket = socket(AF_INET, SOCK_DGRAM)
bind socket to local port number 12000 → serverSocket.bind(("", serverPort))
print ("The server is ready to receive")
loop forever → while True:
    Read from UDP socket into message, getting → message, clientAddress = serverSocket.recvfrom(2048)
    client's address (client IP and port) → modifiedMessage = message.decode().upper()
    send upper case string back to this client → serverSocket.sendto(modifiedMessage.encode(),
    clientAddress)

```

Application Layer: 2-150

150

Socket programming with TCP

Client must contact server

- server process must first be running
- server must have created socket (door) that welcomes client's contact

Client contacts server by:

- Creating TCP socket, specifying IP address, port number of server process
- **when client creates socket:** client TCP establishes connection to server TCP

- when contacted by client, *server TCP creates new socket* for server process to communicate with that particular client
 - allows server to talk with multiple clients
 - *source port numbers used to distinguish clients (more in Chap 3)*

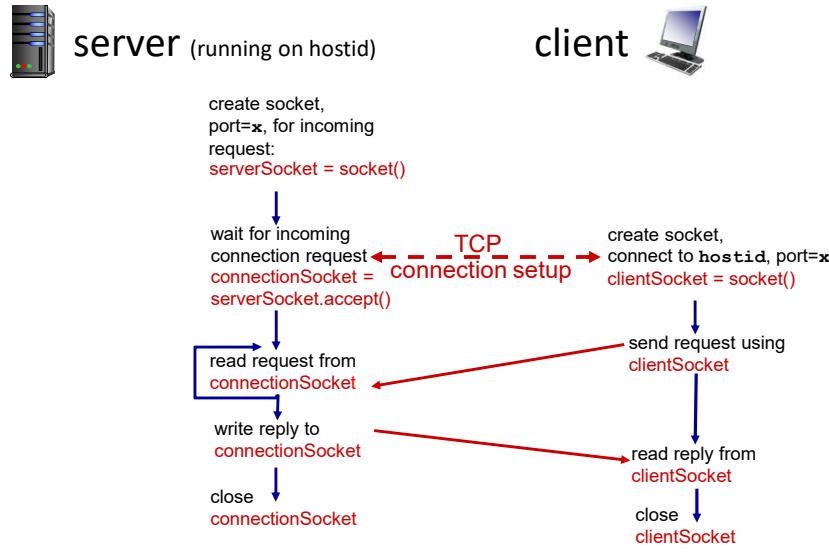
Application viewpoint

TCP provides reliable, in-order byte-stream transfer (“pipe”) between client and server processes

Application Layer: 2-151

151

Client/server socket interaction: TCP



Application Layer: 2-152

152

Example app: TCP client

Python TCPClient

```
from socket import *
serverName = 'servername'
serverPort = 12000
create TCP socket for server,  
remote port 12000 → clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName,serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence.encode())
No need to attach server name, port → modifiedSentence = clientSocket.recv(1024)
print ('From Server:', modifiedSentence.decode())
clientSocket.close()
```

Application Layer: 2-153

153

Example app: TCP server

Python TCPServer

```

from socket import *
serverPort = 12000
create TCP welcoming socket → serverSocket = socket(AF_INET,SOCK_STREAM)
server begins listening for → serverSocket.bind(("",serverPort))
incoming TCP requests → serverSocket.listen(1)
print 'The server is ready to receive'
loop forever → while True:
server waits on accept() for incoming → connectionSocket, addr = serverSocket.accept()
requests, new socket created on return
read bytes from socket (but → sentence = connectionSocket.recv(1024).decode()
not address as in UDP) → capitalizedSentence = sentence.upper()
connectionSocket.send(capitalizedSentence.
encode())
close connection to this client (but not → connectionSocket.close()
welcoming socket)

```

Application Layer: 2-154

154

Chapter 2: Summary

our study of network application layer is now complete!

- application architectures
 - client-server
 - P2P
- application service requirements:
 - reliability, bandwidth, delay
- Internet transport service model
 - connection-oriented, reliable: TCP
 - unreliable, datagrams: UDP
- specific protocols:
 - HTTP
 - SMTP, IMAP
 - DNS
 - P2P: BitTorrent
- video streaming, CDNs
- socket programming:
 - TCP, UDP sockets

Application Layer: 2-155

155

Chapter 2: Summary

Most importantly: learned about *protocols*!

- typical request/reply message exchange:
 - client requests info or service
 - server responds with data, status code
- message formats:
 - *headers*: fields giving info about data
 - *data*: info(payload) being communicated

important themes:

- centralized vs. decentralized
- stateless vs. stateful
- scalability
- reliable vs. unreliable message transfer
- “complexity at network edge”

Application Layer: 2-156