

Main Points

- Process concept

- A process is the OS abstraction for executing a program with limited privileges

- Dual-mode operation: user vs. kernel

- Kernel-mode: execute with complete privileges
- User-mode: execute with fewer privileges

- Safe control transfer

- How do we switch from one mode to the other?

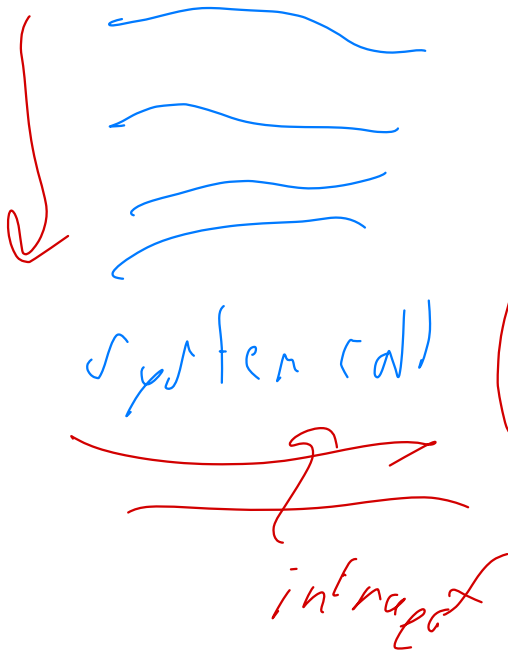
RAM

→ always work

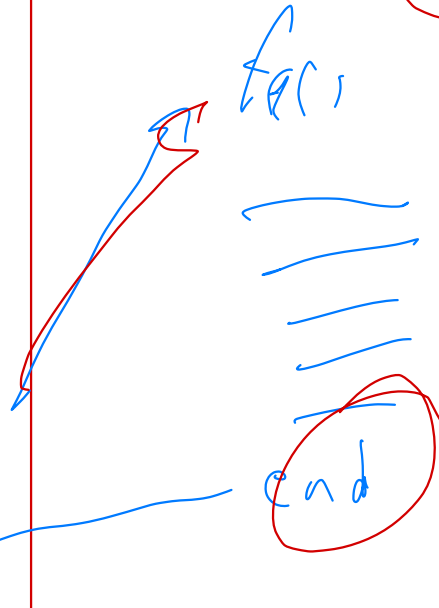
usr

switch mode

kernel



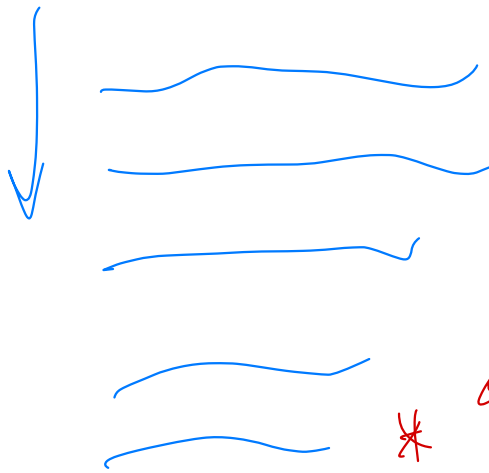
fn



issue n

INT(n)

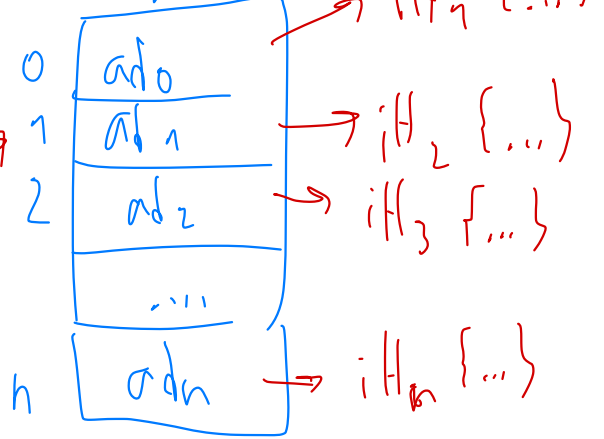
interrupt vector



INT(1)
HW

ih = interrupt handler f

intp vector



Mode Switch

- From user mode to kernel mode
 - Interrupts
 - Triggered by timer and I/O devices
 - Exceptions
 - Triggered by unexpected program behavior
 - Or malicious behavior!
 - System calls (aka protected procedure call)
 - Request by program for kernel to do some operation on its behalf
 - Only limited # of very carefully coded entry points

system call
↑
0
Error R

Mode Switch

- From kernel mode to user mode

- New process/new thread start
 - Jump to first instruction in program/thread

by kernel
switch mode

- Return from interrupt, exception, system call

- Resume suspended execution

- Process/thread context switch

by kernel

- Resume some other process

- User-level upcall (UNIX signal)


- Asynchronous notification to user program

→ user register
by program to kernel to

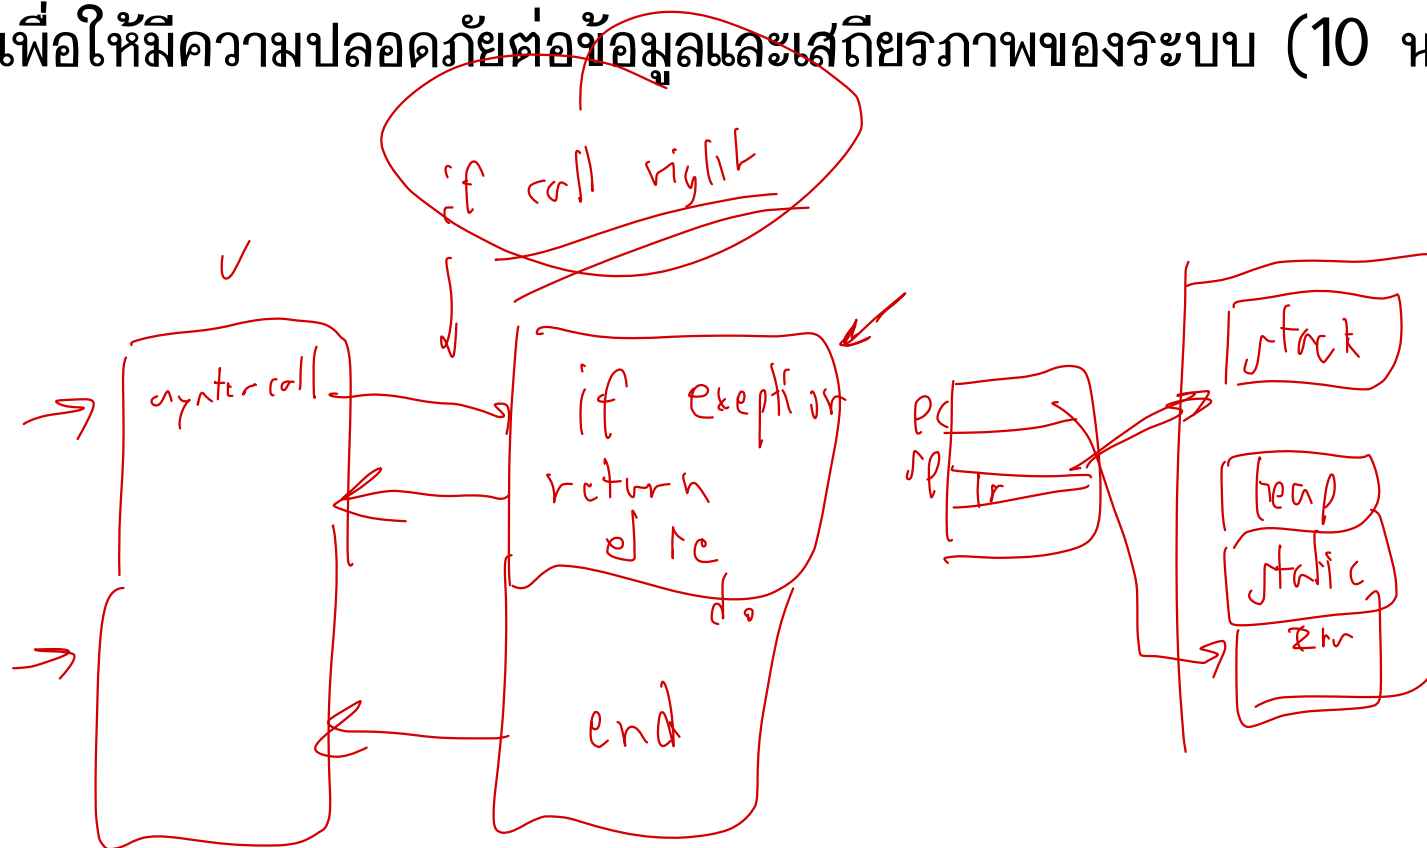
up call

fn → program → user

Activity #1

pc → 

- ในความเห็นของ นศ การทำ mode switch ควรทำอย่างไรบ้าง เพื่อให้มีความปลอดภัยต่อข้อมูลและเสถียรภาพของระบบ (10 นาที)



Implementing Safe Kernel Mode Transfers

- save process state if instruction $\neq 70$
- save register values and save instruction
- must handle weird / buggy / malicious user state
 - system calls with null pointer / call order
 - return instr out of bound / instructions for fail
 - user stack pointer out of bound
- must be able to bug kernel instructions and address
- user must be able to interrupt kernel

Device Interrupts

- OS kernel needs to communicate with physical devices

- Devices operate asynchronously from the CPU

- Polling: Kernel waits until I/O is done

- Interrupts: Kernel can do other work in the meantime

- Device access to memory

- Programmed I/O: CPU reads and writes to device

- Direct memory access (DMA) by device

- Buffer descriptor: sequence of DMA's

- E.g., packet header and packet body

- Queue of buffer descriptors

- Buffer descriptor itself is DMA'ed

1 descriptor = 1 op

Program → File
RAM

học tập là int
bổn m f

26
lưu trữ
đang chờ
đợi

lưu trữ
đang chờ
đợi

lưu trữ
đang chờ
đợi

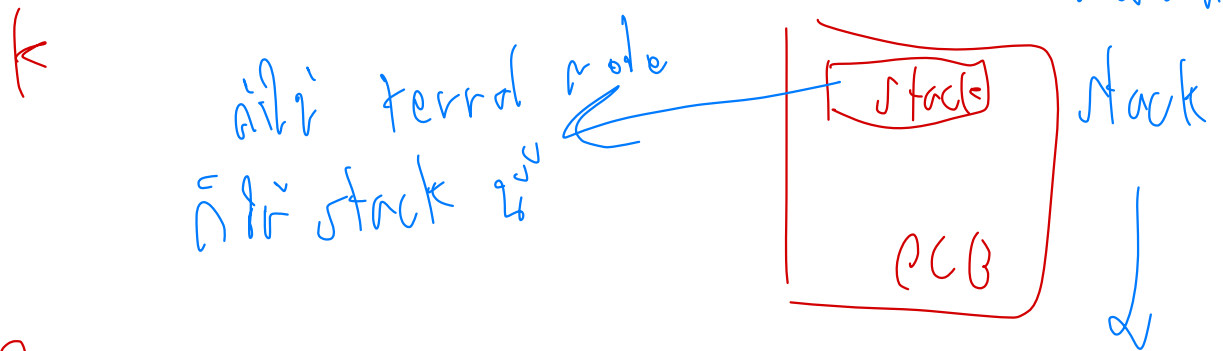
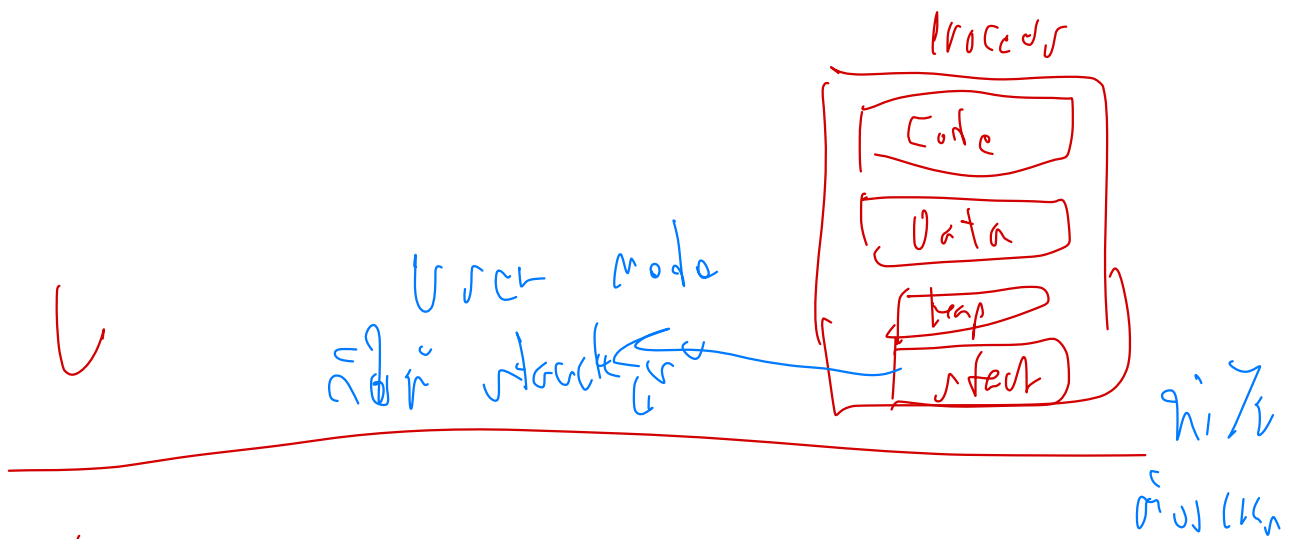
lưu trữ
đang chờ
đợi

Activity #2

- How do device interrupts work?

- Where does the CPU run after an interrupt? ถ้า interrupt เกิด
- What stack does it use? → ถ้า interrupt มีโปรแกรม ให้ใช้ stack
- Is the work the CPU had been doing before the process ที่ interrupt lost forever? ใช่แล้ว 70 ถ้า int ใน kernel
- If not, how does the CPU know how to resume that work? ถ้า ดูจาก register ใน stack

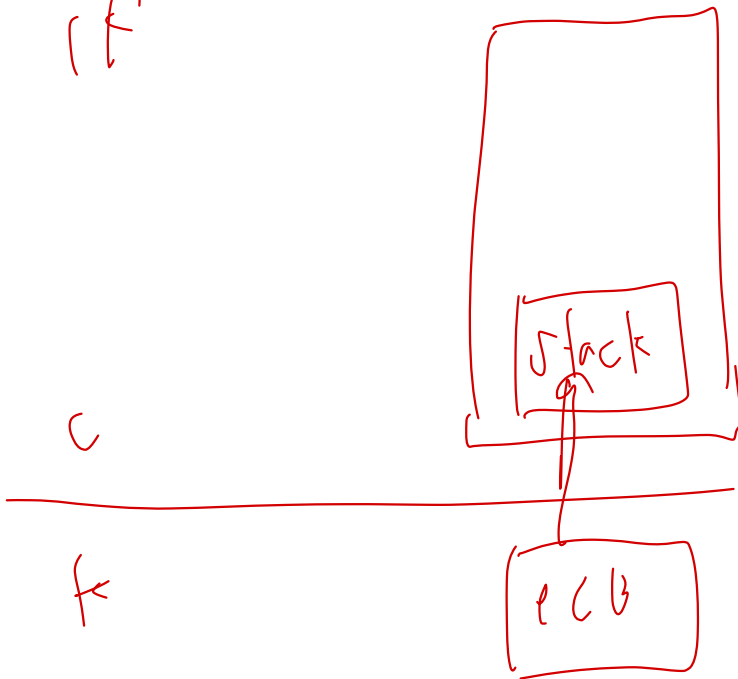
(10 นาที)



if

address kernel

security



use user API

int native distribution

system call process

kernel stack

kernel

kernel stack

kernel

slide 22

pp 02

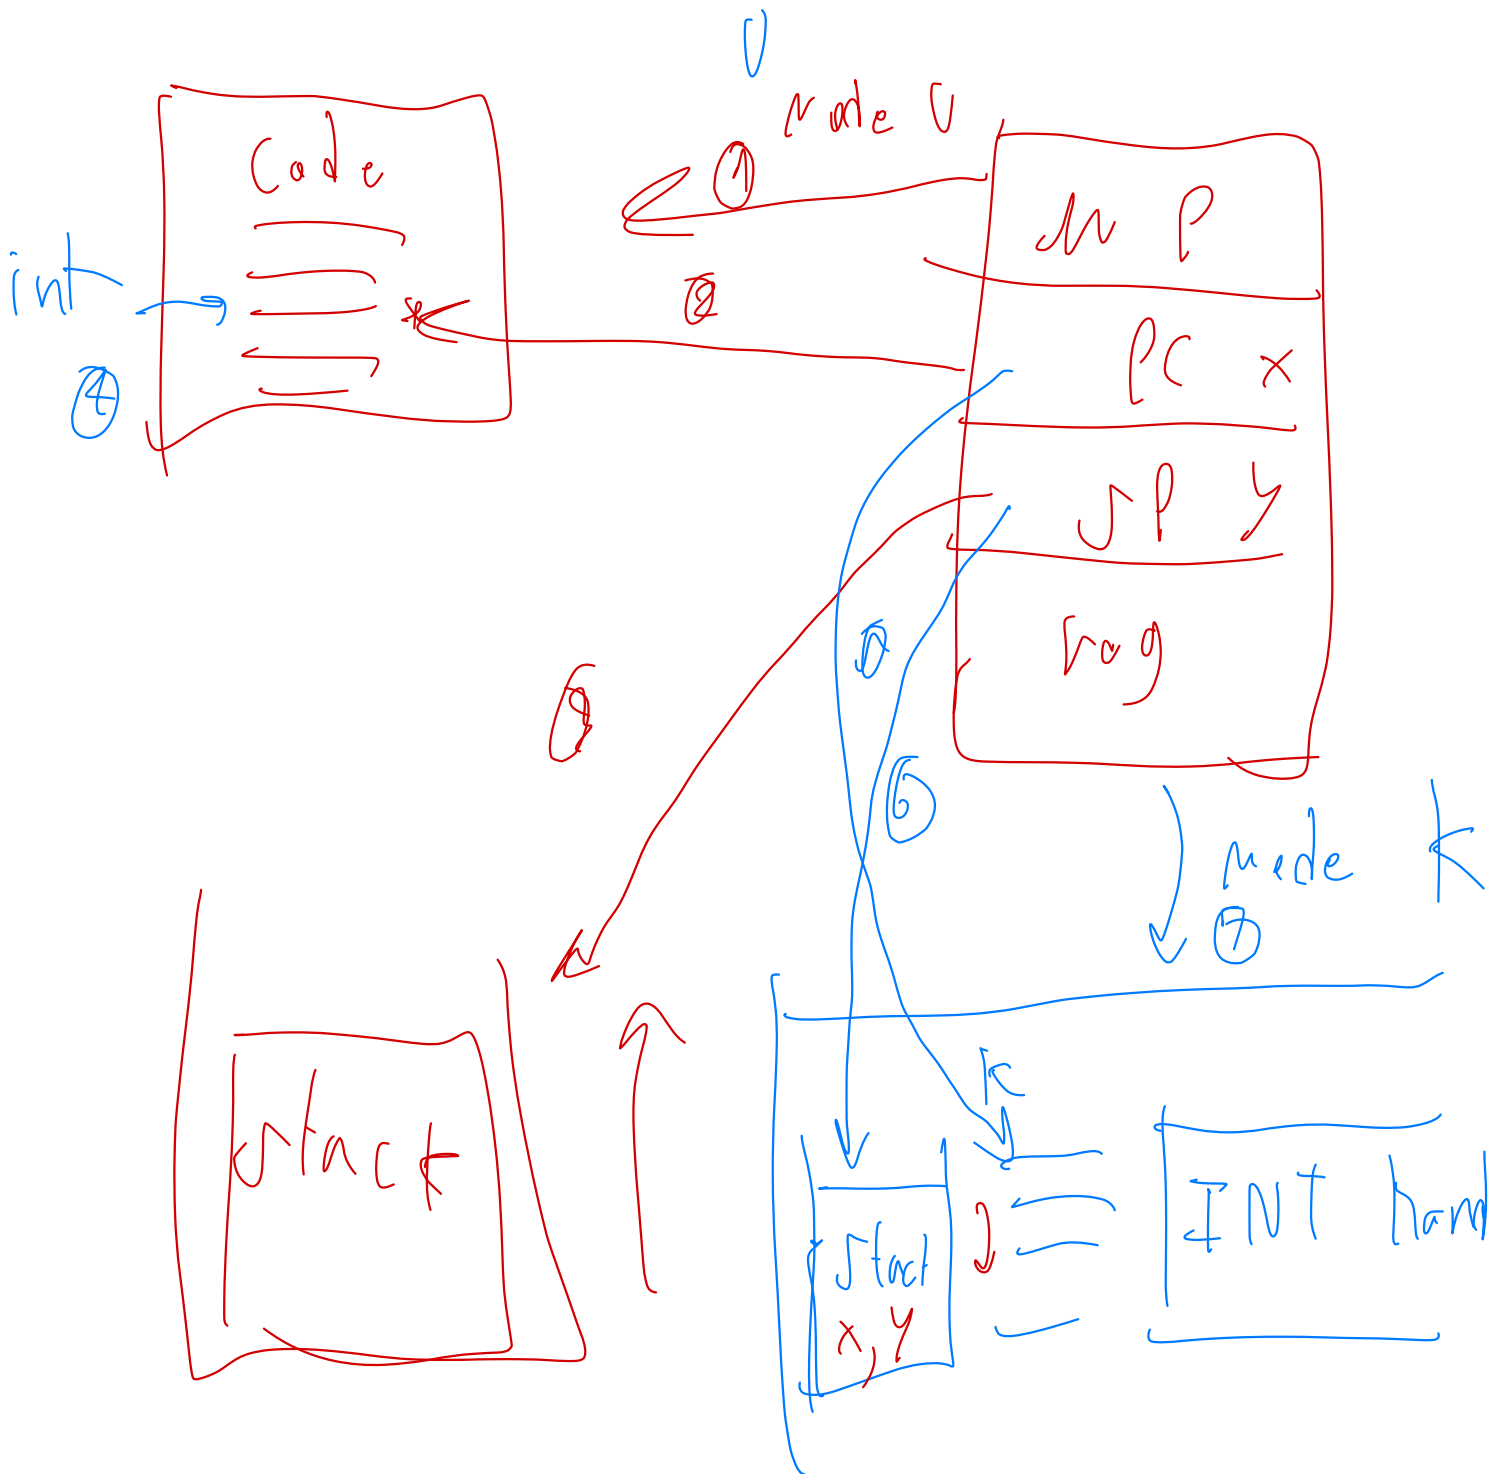
How do we take interrupts safely?

- Interrupt vector
 - Limited number of entry points into kernel
- Kernel interrupt stack
 - Handler works regardless of state of user code
- Interrupt masking → *beginning int vs beginning int rule*
 - Handler is non-blocking
- Atomic transfer of control
 - "Single instruction"-like to change:
 - Program counter → *PC*
 - Stack pointer → *SP*
 - Memory protection → *MP*
 - Kernel/user mode → *KU*
- Transparent restartable execution
 - User program does not know interrupt occurred

(OS (system))

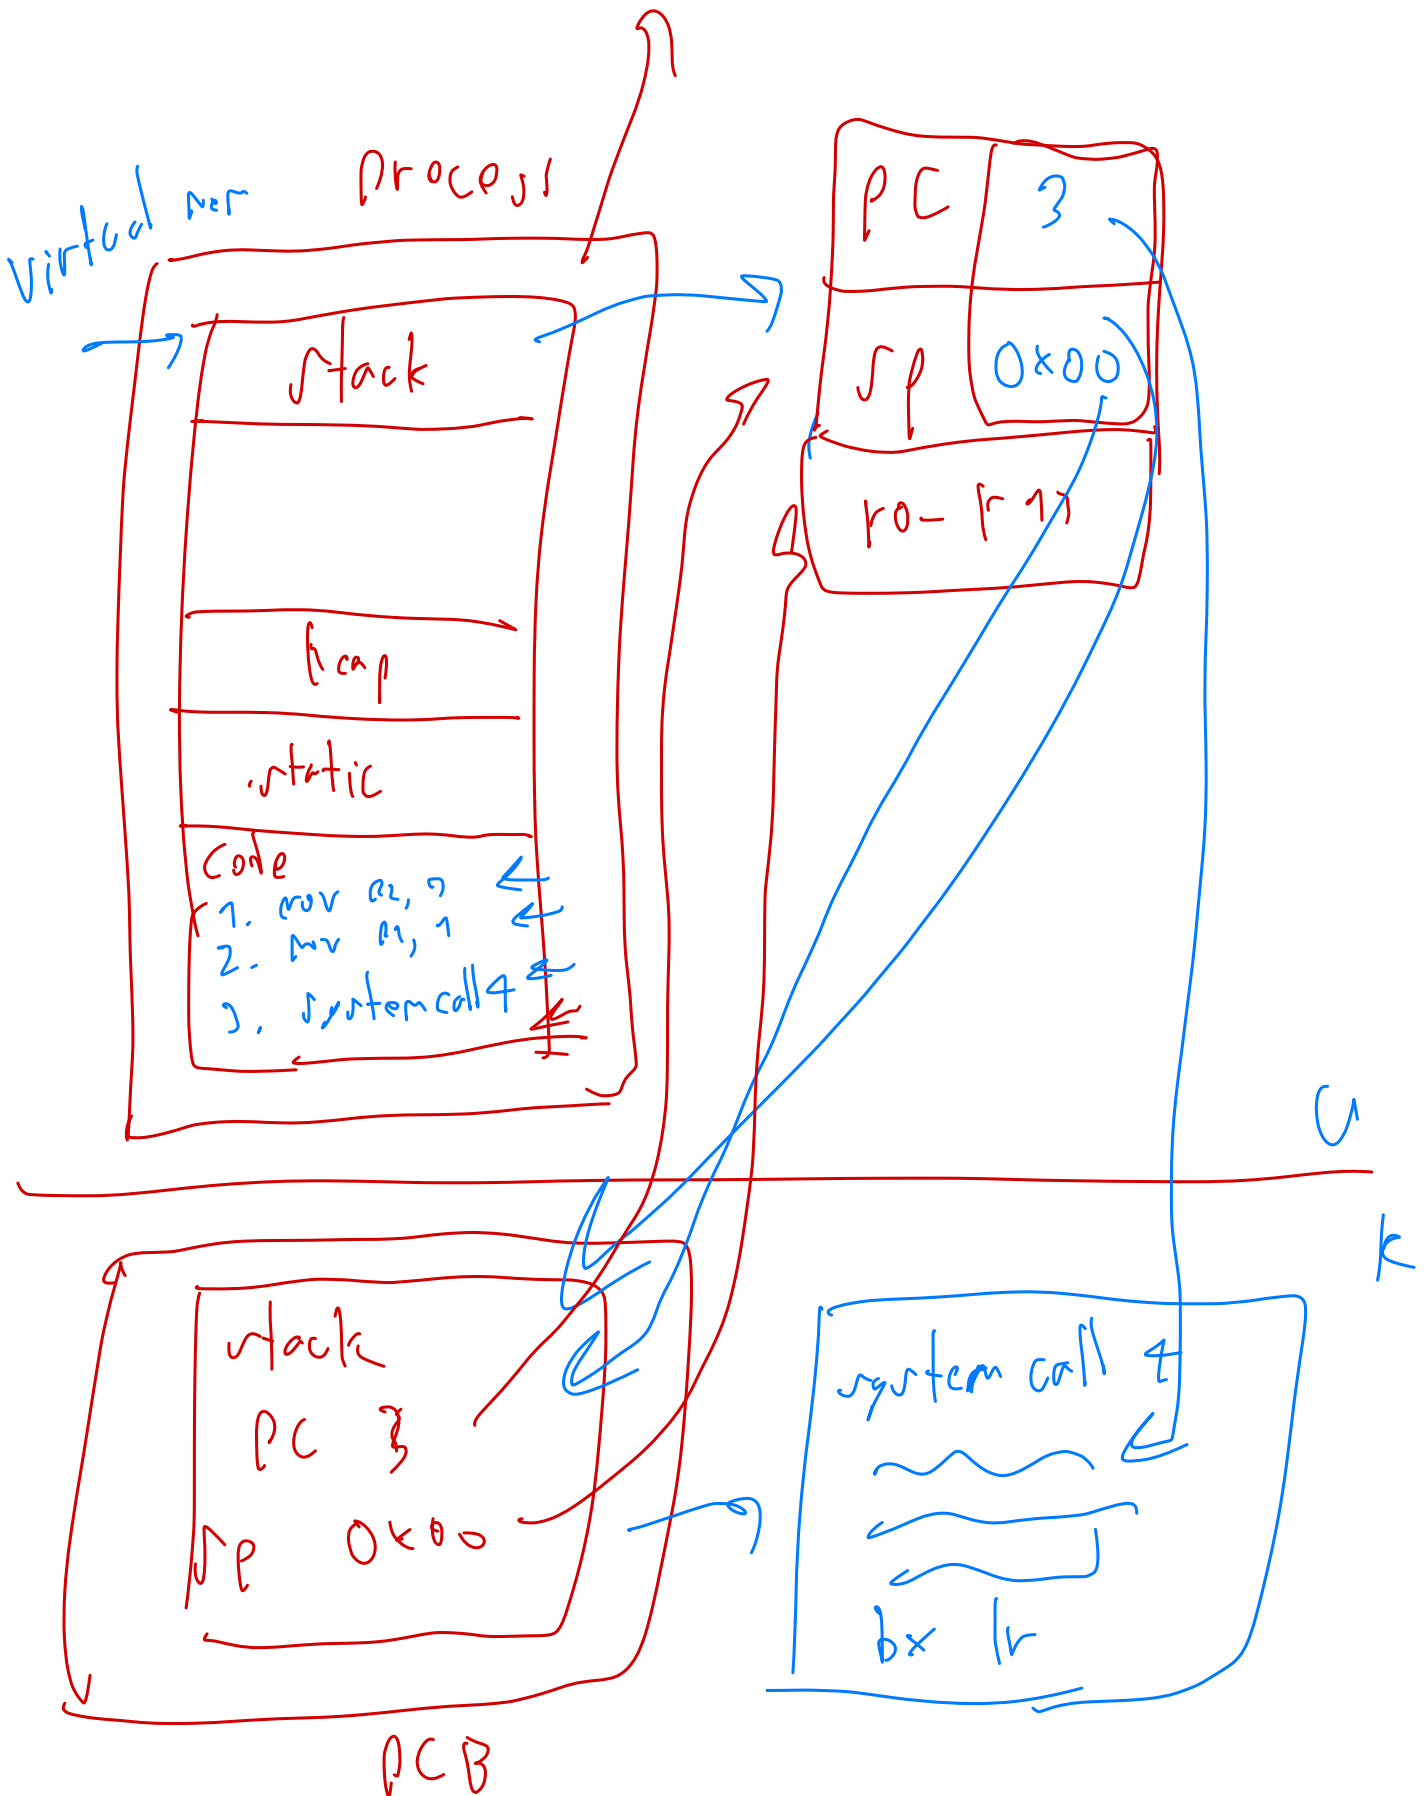
1 instr beginning of instr

↑ instr



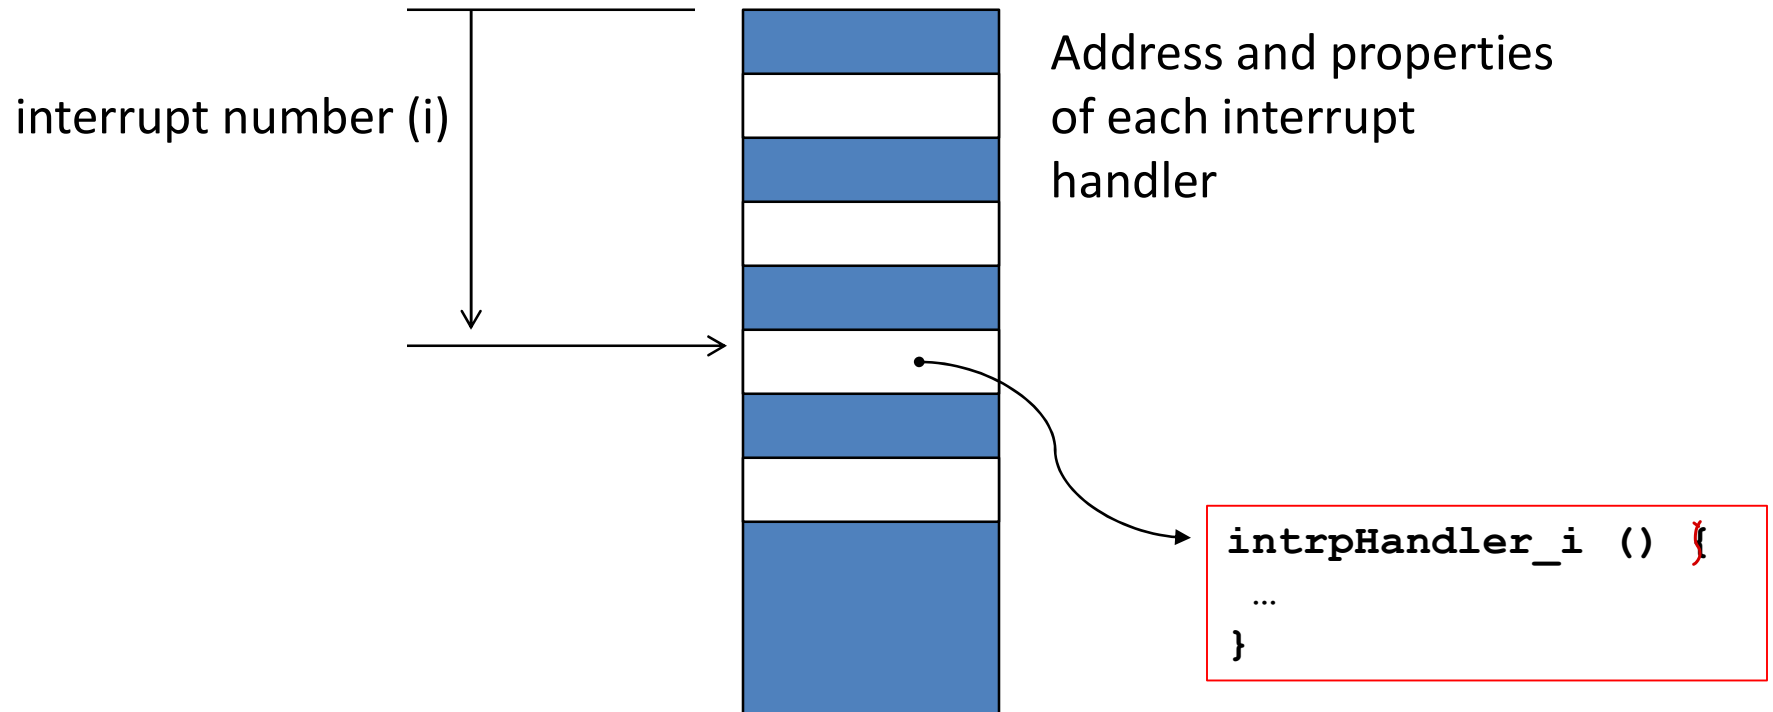
మూలం నాది x, y వీలు నాది
 ఓకే pointer నాది switch mode

Process



Where do mode transfers go?

- Solution: ***Interrupt Vector***



The Kernel Stack

- Interrupt handlers want a stack
- System call handlers want a stack
- Can't just use the user stack [why?]

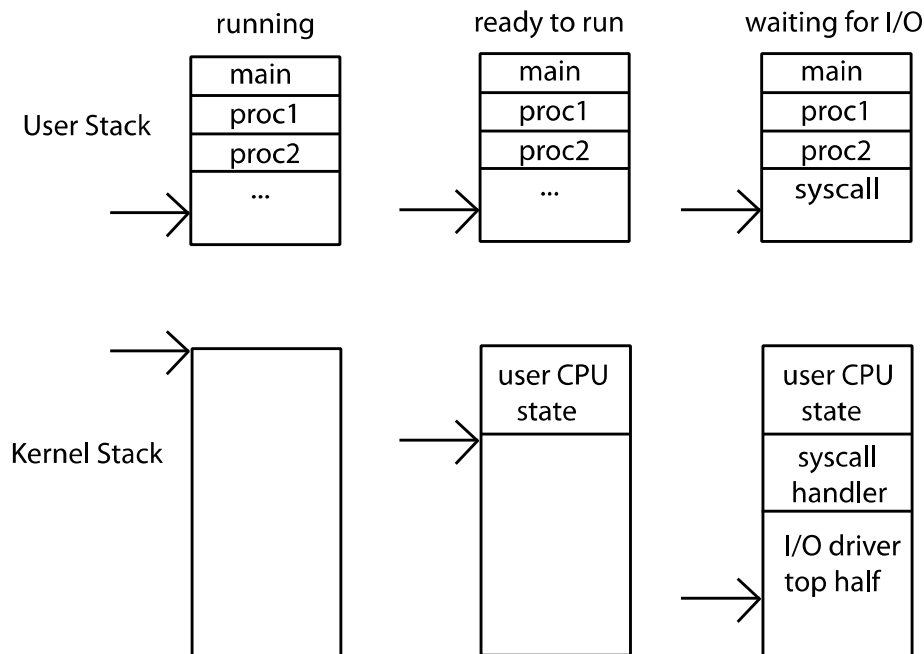
no! user

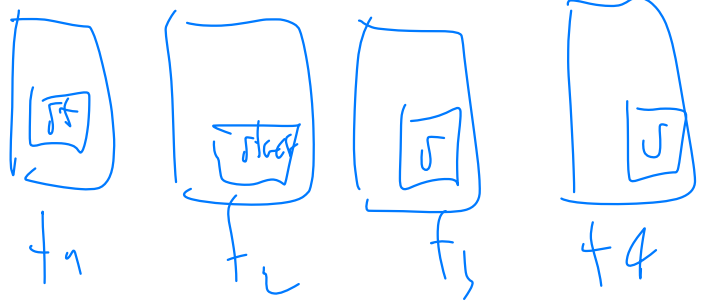
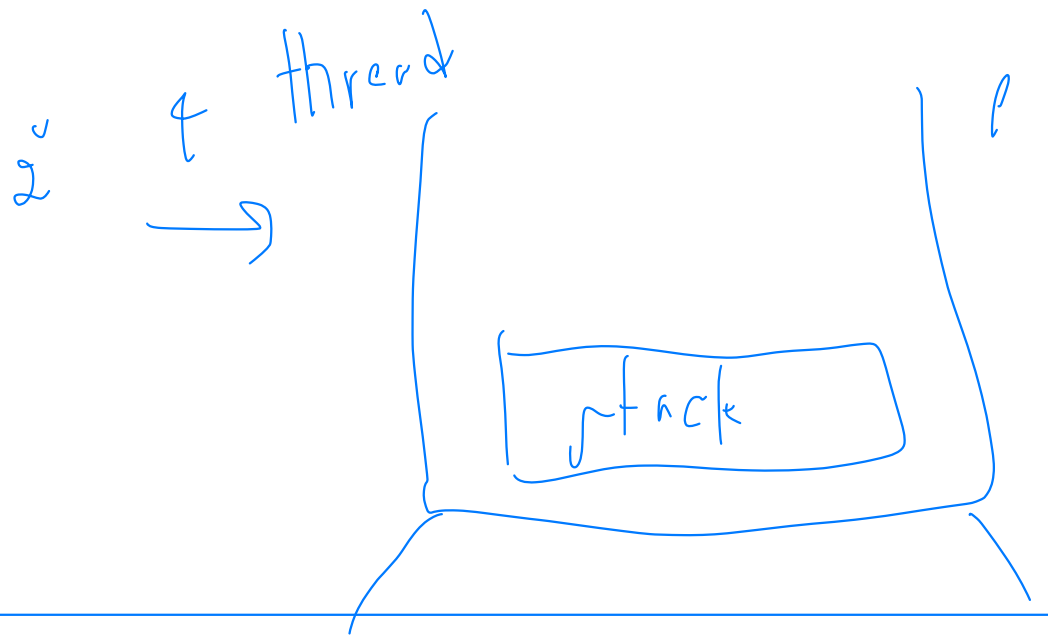
provide multi-thread
like window

process \rightarrow n thread

The Kernel Stack

- Solution: two-stack model
 - Each OS thread has kernel stack (located in kernel memory) plus user stack (located in user memory)
- Place to save user registers during interrupt





Interrupt Stack

- Per-processor, located in kernel (not user) memory
 - Usually a process/thread has both: kernel and user stack
- Why can't the interrupt handler run on the stack of the interrupted user process?

Case Study: x86 Interrupt

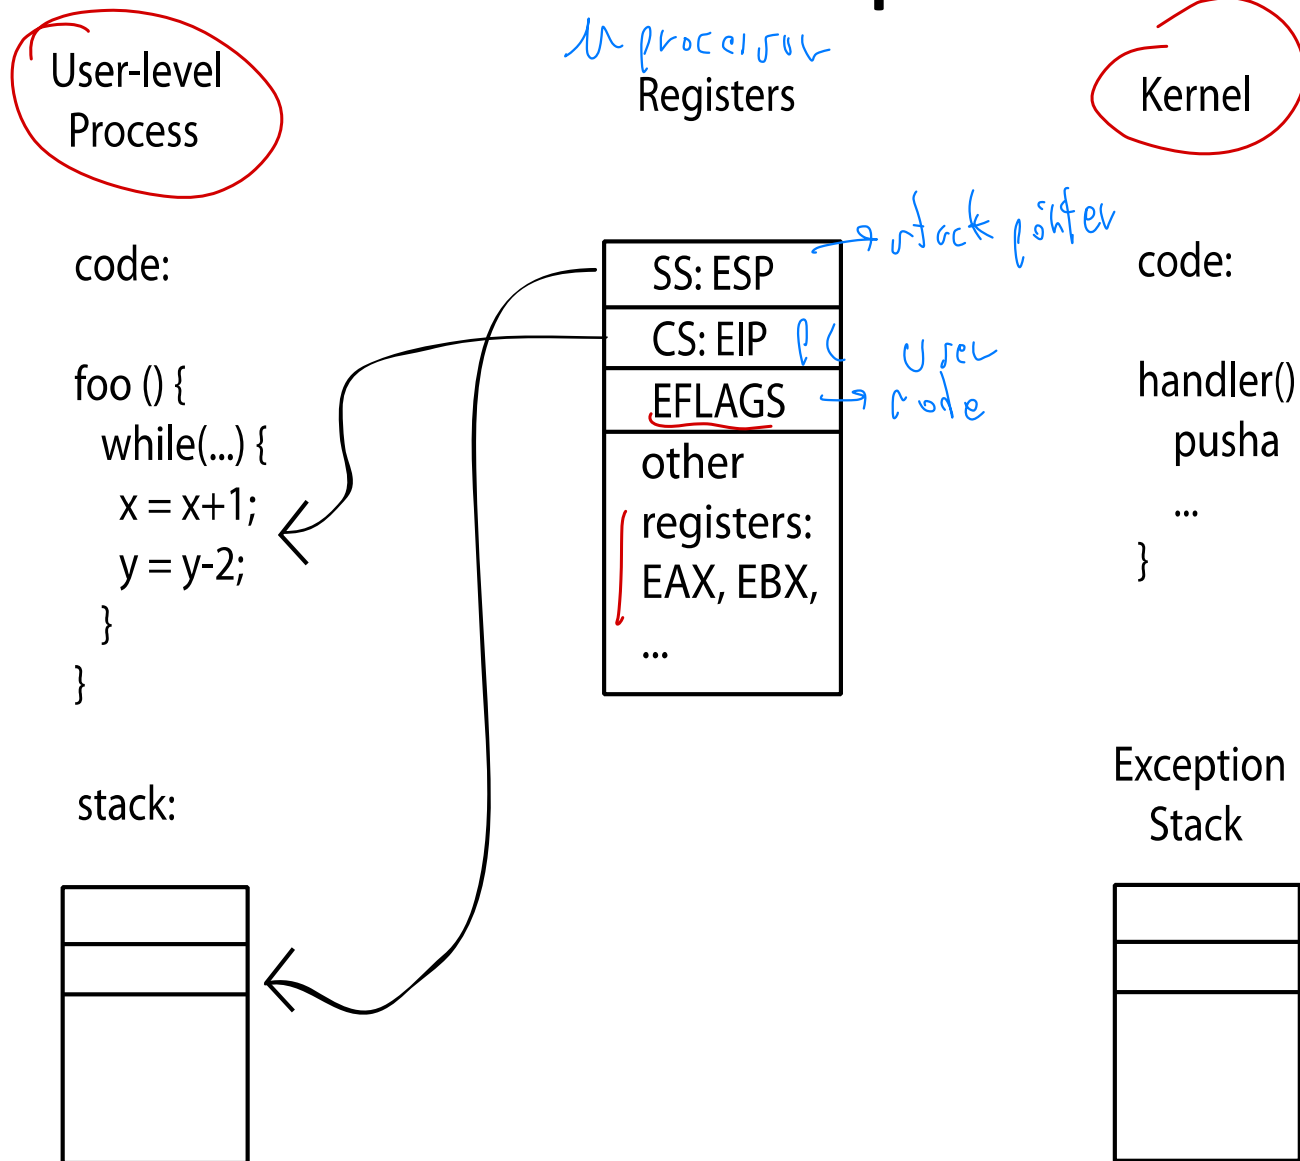
- Save current stack pointer
- Save current program counter
- Save current processor status word (condition codes)
- Switch to kernel stack; put SP, PC, PSW on stack
- Switch to kernel mode
- Vector through interrupt table
- Interrupt handler saves registers it might clobber

save
auto

save

→ in routine *initializing*
intcs *no* *etc*

Before Interrupt



During Interrupt

User-level
Process

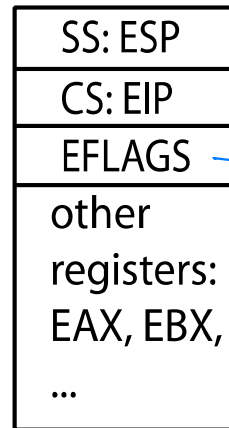
code:

```
foo () {  
    while(...) {  
        x = x+1;  
        y = y-2;  
    }  
}
```

stack:



Registers

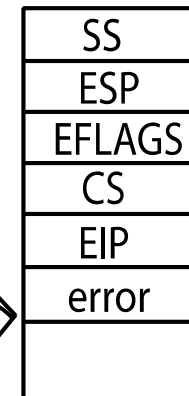


Kernel

code:

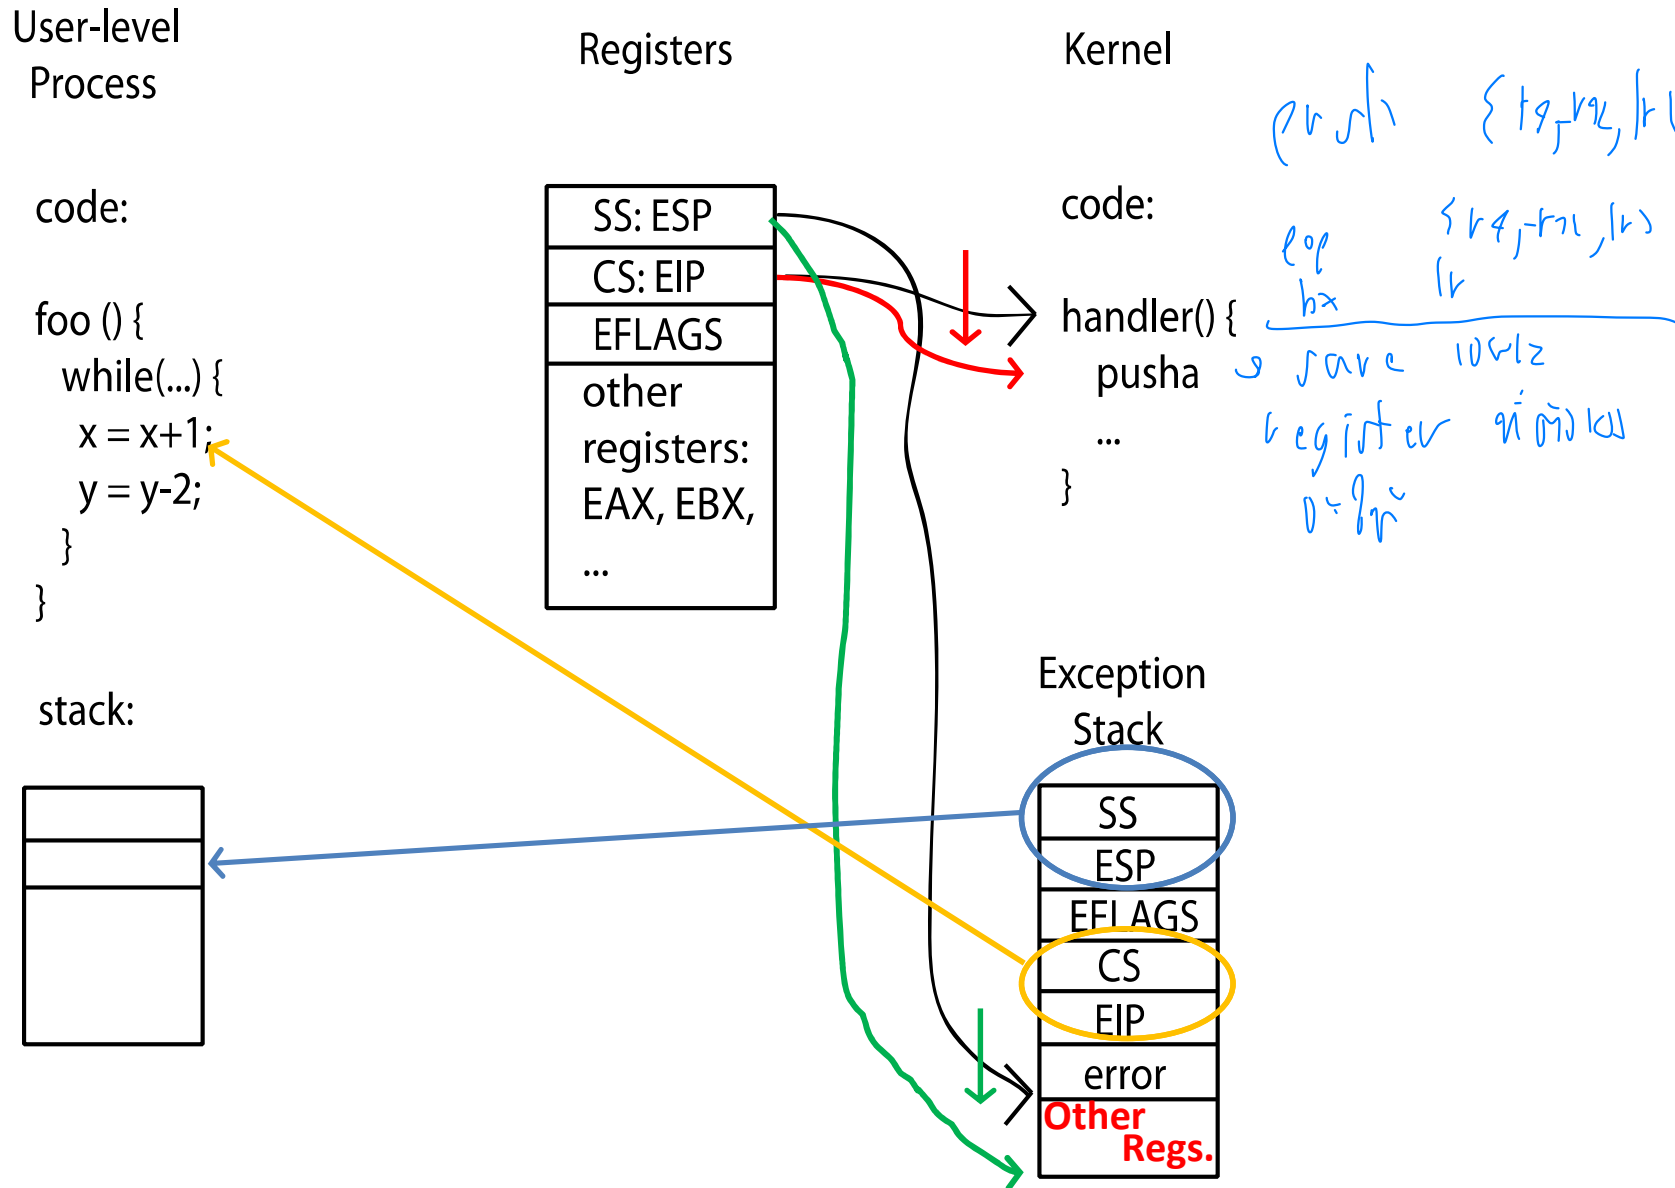
```
handler() {  
    pusha  
    ...  
}
```

Exception
Stack



} mp
} PC
⊗ int uia S
information n
tag 21000

After Interrupt



At end of handler

- Handler restores saved registers
- Atomically return to interrupted process/thread
 - Restore program counter
 - Restore program stack
 - Restore processor status word/condition codes
 - Switch to user mode

Interrupt Masking

- Interrupt handler runs with interrupts off *since int not desirable vs int by cpu*
 - Re-enabled when interrupt completes
- OS kernel can also turn interrupts off
 - Eg., when determining the next process/thread to run
 - On x86
 - CLI: disable interrupts
 - STI: enable interrupts
 - Only applies to the current CPU (on a multicore)
- We'll need this to implement synchronization in chapter 5

interruptible int qds M processor

Hardware support: Interrupt Control

- Interrupt processing not visible to the user process:
 - Occurs between instructions, restarted transparently
 - No change to process state
 - What can be observed even with perfect interrupt processing?

int qds int qds int qds

- Interrupt Handler invoked with interrupts 'disabled'
 - Re-enabled upon completion
 - Non-blocking (run to completion, no waits)
 - Pack up in a queue and pass off to an OS thread for hard work

- wake up an existing OS thread

Thread | 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100

Hardware support: Interrupt Control

- OS kernel may enable/disable interrupts
 - On x86: CLI (disable interrupts), STI (enable)
 - Atomic section when select next process/thread to run
 - Atomic return from interrupt or syscall

- HW may have multiple levels of interrupts
 - Mask off (disable) certain interrupts, eg., lower priority
 - Certain Non-Maskable-Interrupts (NMI)
 - e.g., kernel segmentation fault
 - Also: Power about to fail!

int
segmentation fault

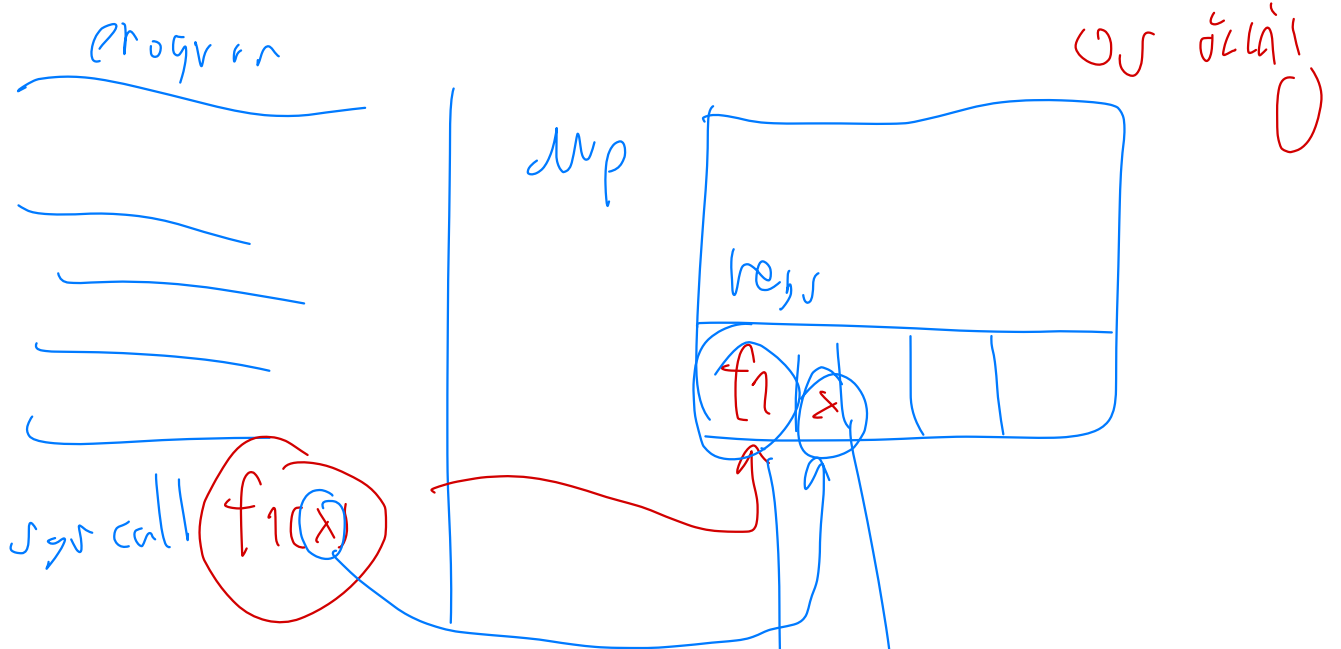
power

int 0
int 1
int level 0
int level 1
prop 541010

Doing the syscall inside int

Kernel System Call Handler

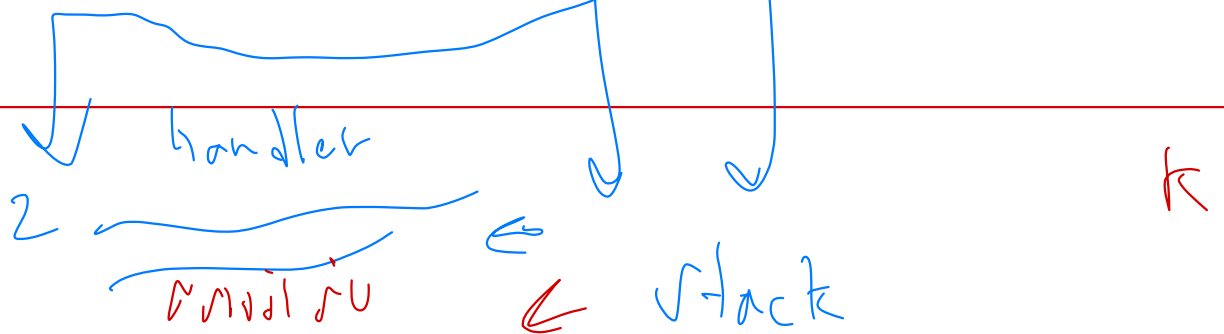
- **Vector through well-defined syscall entry points!**
 - Table mapping system call number to handler
- Locate arguments
 - In registers or on user (!) stack
- Copy arguments
 - From user memory into kernel memory – carefully checking locations!
 - Protect kernel from malicious code evading checks
- Validate arguments
 - Protect kernel from errors in user code
- Copy results back
 - Into user memory – carefully checking locations!



cauza eșec code nime = **trap**

nivelul în stivă int

la 0 în stivă int lui f1 în int reală



parametri

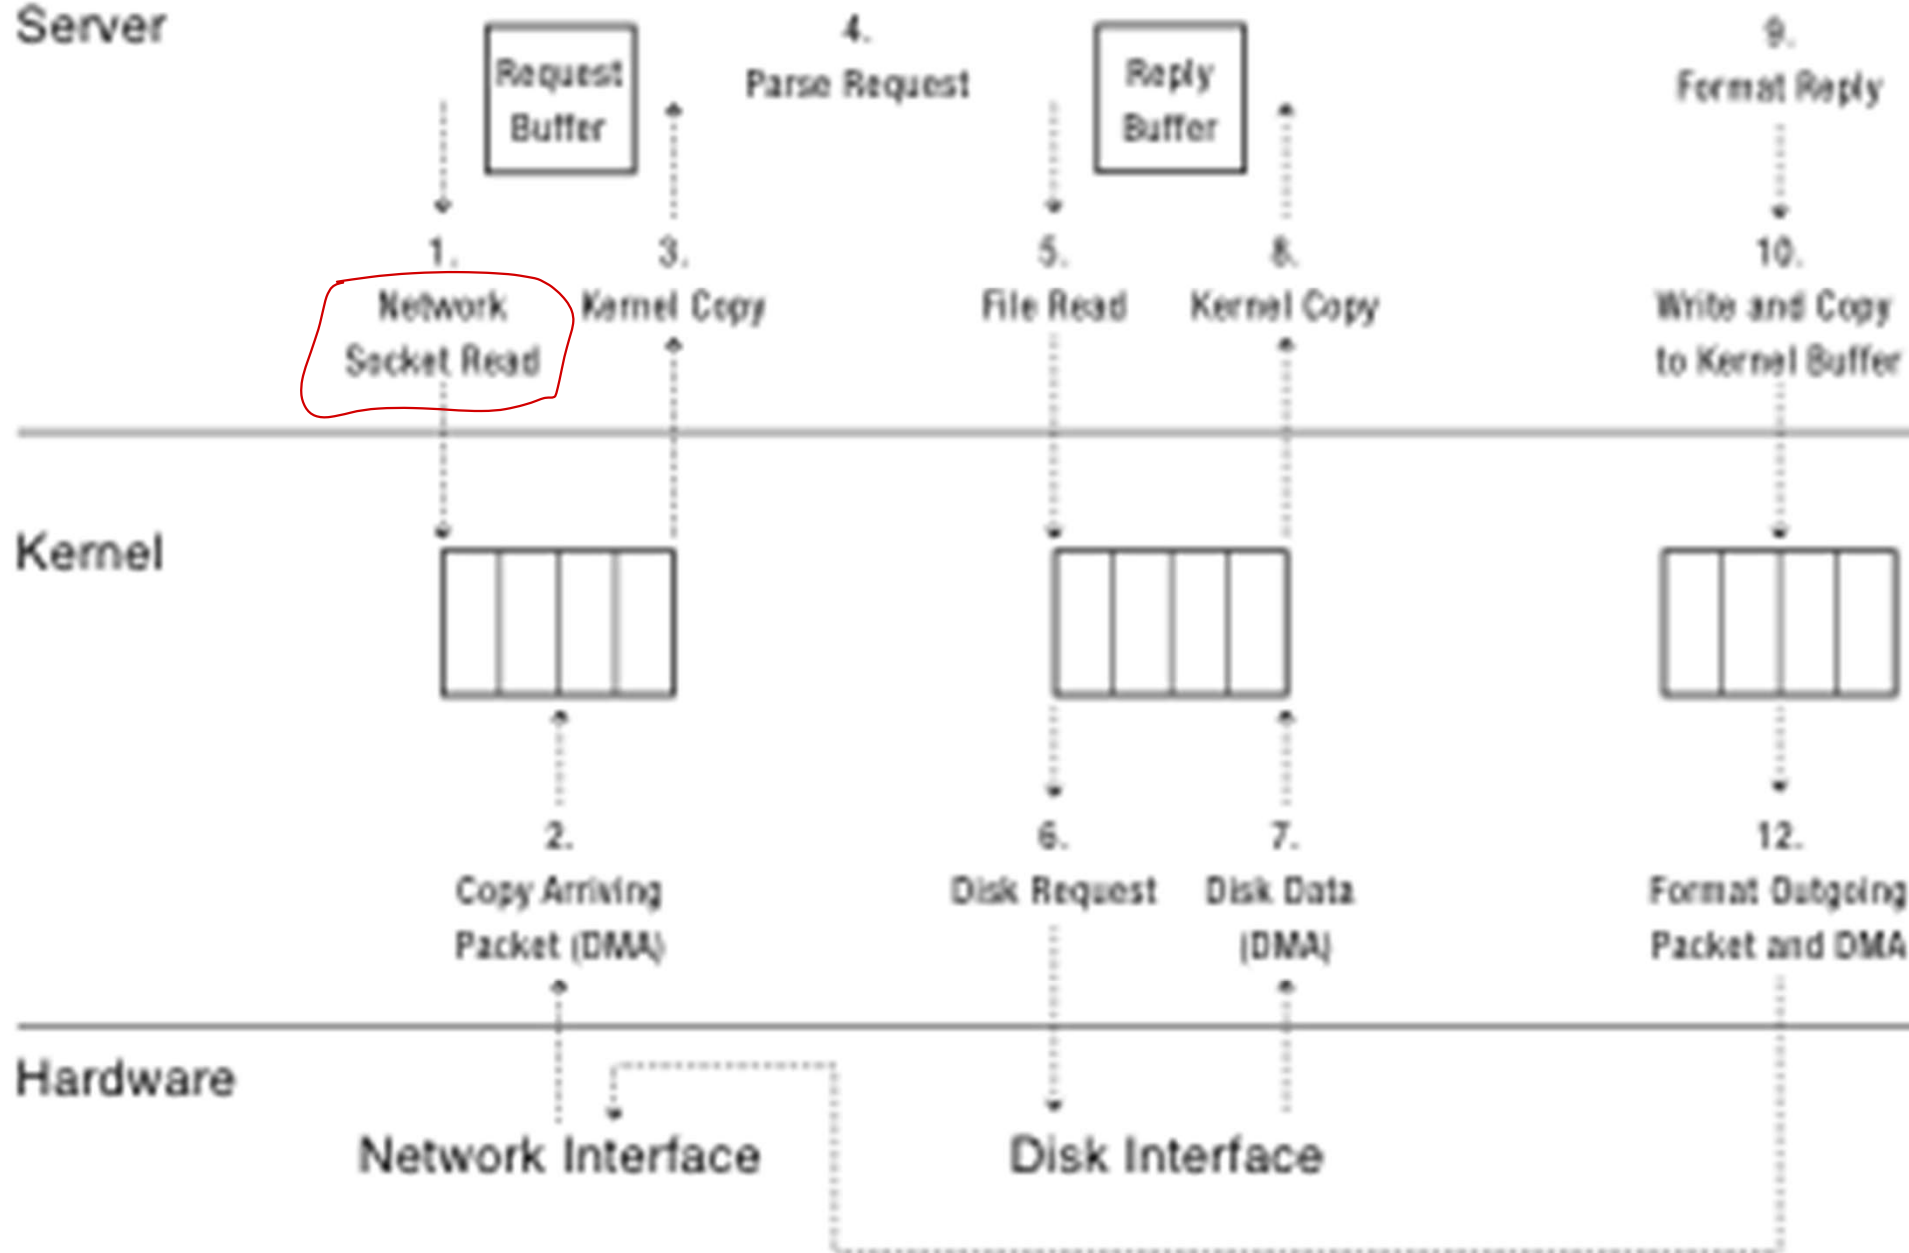
argumente

syscall f1(x)

la x86 trap în stivă în syscall
în stivă în Trap

nvme network socket read

Server



Today: Four Fundamental OS Concepts

- **Thread: Execution Context** *vir Processor*
 - Program Counter, Registers, Execution Flags, Stack
- **Address space (with translation)** *virtual*
sep *← real*
all mem from
 - Program's view of memory is distinct from physical machine
- **Process: an instance of a running program**
 - Address Space + One or more Threads
- **Dual mode operation / Protection** *U/K* *ni lĩ cĩn*
 - Only the “system” can access certain resources
 - Combined with translation, isolates programs from each other

protect all vi ĩs Hw,

ĩĩ kernel run nĩĩĩ